



**Hewlett Packard**  
Enterprise

# **Using Flex Tables**

HPE Vertica Analytic Database

Software Version: 7.2.x

**Document Release Date: 12/7/2015**

## **Legal Notices**

### **Warranty**

The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. HPE shall not be liable for technical or editorial errors or omissions contained herein.

The information contained herein is subject to change without notice.

### **Restricted Rights Legend**

Confidential computer software. Valid license from HPE required for possession, use or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

### **Copyright Notice**

© Copyright 2015 Hewlett Packard Enterprise Development LP

### **Trademark Notices**

Adobe™ is a trademark of Adobe Systems Incorporated.

Microsoft® and Windows® are U.S. registered trademarks of Microsoft Corporation.

UNIX® is a registered trademark of The Open Group.

This product includes an interface of the 'zlib' general purpose compression library, which is Copyright © 1995-2002 Jean-loup Gailly and Mark Adler.

# Contents

Getting Started .....	8
Create a Simple JSON File .....	8
Create a Flex Table and Load Data .....	8
Query More of Your Flex Table .....	9
Build a Flex Table View .....	11
Create a Hybrid Flex Table .....	12
Promote Virtual Columns in a Hybrid Flex Table .....	13
Understanding Flex Tables .....	17
Exploration to Promotion .....	17
Flex Table Terms .....	18
Is There Structure in a Flex Table? .....	18
Making Flex Table Data Persist .....	18
What Happens When You Create Flex Tables? .....	19
Creating Superprojections Automatically .....	21
Default Flex Table View .....	21
Flex Functions .....	22
Using Clients with Flex Tables .....	22
Creating Flex Tables .....	23
Unsupported CREATE FLEX TABLE Statements .....	23
Creating Basic Flex Tables .....	23
Creating Temporary Flex Tables .....	23
Materializing Flex Table Virtual Columns .....	25
Creating Columnar Tables from Flex Tables .....	26
Creating External Flex Tables .....	27
Partitioning Flex Tables .....	28
Using COPY with Flex Tables .....	29
Basic Flex Table Load and Query .....	30
Loading Data into Flex Table Real Columns .....	30
Handling Default Values During Loading .....	32
Using COPY to Specify Default Column Values .....	33

Using Flex Tables for IDOL Data .....	34
ODBC Connection String for CFS .....	35
CFS COPY LOCAL Statement .....	35
Using Flex Table Parsers .....	36
Using Flex Parsers for Columnar Tables .....	36
Loading Avro Data .....	38
Rejecting Data on Materialized Column Type Errors .....	39
Primitive Data Types for favroparser .....	39
Complex Data Types for favroparser .....	40
Records .....	41
Enums .....	42
Arrays .....	42
Maps .....	43
Unions .....	44
Fixed .....	44
Loading Common Event Format (CEF) Data .....	44
Create a Flex Table and Load CEF Data .....	45
Create a Columnar Table and Load CEF Data .....	46
Compute Keys and Build a Flex Table View .....	47
Use the fcefpaser Delimiter Parameter .....	48
Loading CSV Data .....	49
Using Default Parser Settings .....	49
Rejecting on Duplicate Values .....	52
Rejecting Data on Materialized Column Type Errors .....	53
Rejecting or Omitting Empty Rows .....	55
Loading Delimited Data .....	58
Rejecting Duplicate Values .....	58
Rejecting Materialized Column Type Errors .....	59
Loading JSON Data .....	61
Checking JSON Integrity .....	61
Using flatten_maps and flatten_arrays Parameters .....	61
Loading from a Specific Start Point .....	62
Parsing From a Start Point Occurrence .....	63
Controlling Column Name Separators .....	64
Handling Special Characters .....	64
Rejecting on Duplicate Values .....	64
Rejecting Data on Materialized Column Type Errors .....	64
Rejecting or Omitting Empty Rows .....	66
Loading Matches from Regular Expressions .....	67
Sample Regular Expression .....	67
Using Regular Expression Matches for a Flex Table .....	67
Using fregexparser for Columnar Tables .....	67
Using External Tables with fregexparser .....	68

Computing Flex Table Keys .....	70
Using COMPUTE_FLEXTABLE_KEYS .....	70
Calculating Key Value Column Widths .....	70
Materializing Flex Tables .....	72
Adding Columns to Flex Tables .....	72
Adding Columns with Default Values .....	73
Changing the __raw__ Column Size .....	75
Changing Flex Table Real Columns .....	76
Dropping Flex Table Columns .....	76
Updating Flex Table Views .....	77
Using BUILD_FLEXTABLE_VIEW .....	77
Handling JSON Duplicate Key Names in Views .....	78
Creating a Flex Table View .....	80
Using COMPUTE_FLEXTABLE_KEYS_AND_BUILD_VIEW .....	82
Querying Flex Tables .....	83
Unsupported DDL and DML Statements .....	83
Querying Flex Table Keys .....	83
Determining Flex Table Data Contents .....	84
Querying Virtual Columns .....	84
Using Functions and Casting in Flex Table Queries .....	85
Casting Data Types in a Query .....	86
Accessing an Epoch Key .....	86
Querying Flex Views .....	87
Listing Flex Tables .....	88
Setting Flex Table Parameters .....	89
Flex Data Functions Reference .....	91
Flex Table Dependencies .....	92
Associating Flex Tables and Views .....	92
BUILD_FLEXTABLE_VIEW .....	92
COMPUTE_FLEXTABLE_KEYS .....	96

COMPUTE_FLEXTABLE_KEYS_AND_BUILD_VIEW .....	98
MATERIALIZED_FLEXTABLE_COLUMNS .....	98
RESTORE_FLEXTABLE_DEFAULT_KEYS_TABLE_AND_VIEW .....	101
<b>Flex Extractor Functions Reference .....</b>	<b>103</b>
MAPDELIMITEDEXTRACTOR .....	103
Parameters .....	103
Examples .....	104
See Also .....	106
MAPJSONEXTRACTOR .....	107
Parameters .....	107
Examples .....	108
See Also .....	110
MAPREGEXEXTRACTOR .....	110
Parameters .....	110
Examples .....	110
See Also .....	111
<b>Flex Map Functions Reference .....</b>	<b>113</b>
EMPTYMAP .....	113
MAPAGGREGATE .....	115
MAPCONTAINSKEY .....	116
MAPCONTAINSVALUE .....	118
MAPITEMS .....	120
MAPKEYS .....	123
MAPKEYSINFO .....	125
MAPLOOKUP .....	127
MAPSIZE .....	134
MAPTOSTRING .....	135
MAPVALUES .....	138
MAPVERSION .....	139
<b>Flex Parsers Reference .....</b>	<b>141</b>
FAVROPARSER .....	141
Parameters .....	141
Examples .....	142
See Also .....	143
FCEFPARSER .....	143
Parameters .....	143
Examples .....	144

See Also .....	145
FCSVPARSER .....	145
Parameters .....	145
Examples .....	147
See Also .....	148
FDELIMITEDPAIRPARSER .....	148
Parameters .....	148
Examples .....	149
See Also .....	151
FDELIMITEDPARSER .....	151
Parameters .....	151
Examples .....	153
See Also .....	153
FJSONPARSER .....	154
Parameters .....	154
Examples .....	155
FREGEXPARSER .....	156
Parameters .....	156
Example .....	157
See Also .....	159
Send Documentation Feedback .....	161

# Getting Started

Getting Started describes the basics of creating, exploring, and using flex tables. The rest of this guide presents *beyond the basics* details using simple examples.

## Create a Simple JSON File

Use this JSON data for the exercises in the rest of this section:

```
{"name": "Everest", "type": "mountain", "height": 29029, "hike_safety": 34.1}
{"name": "Mt St Helens", "type": "volcano", "height": 29029, "hike_safety": 15.4}
{"name": "Denali", "type": "mountain", "height": 17000, "hike_safety": 12.2}
{"name": "Kilimanjaro", "type": "mountain", "height": 14000 }
{"name": "Mt Washington", "type": "mountain", "hike_safety": 50.6}
```

1. Copy and paste the JSON data into your favorite editor.
2. Save the file in any convenient location for loading into your Vertica database.

## Create a Flex Table and Load Data

1. Create a flex table called `mountains`:

```
=> CREATE flex table mountains();
```

2. Load the JSON file you saved, using the flex table parser `fjsonparser`:

```
=> COPY mountains from '/home/dbadmin/data/flex/mountains.json' parser fjsonparser();
Rows Loaded
-----
                5
(1 row)
```

3. Query values from the sample file:

```
=> SELECT name, type, height from mountains;
  name      | type    | height
-----+-----+-----
Everest     | mountain | 29029
Mt St Helens | volcano  | 29029
Denali      | mountain | 17000
Kilimanjaro | mountain | 14000
Mt Washington | mountain |
(5 rows)
```



You have now created a flex table and loaded data. Next, learn more about using flex table data in your database.

## Query More of Your Flex Table

1. Query your flex table to see the data you loaded as it is stored in the `__raw__` column. The example illustrates the table contents, with Return characters added for illustration:

```
=> \x
Expanded display is on.
=> SELECT * from mountains;

[ RECORD 1 ]+-----
__identity__ | 1
__raw__      | \001\000\000\000,\000\000\000\004\000\000\000\024\000\000\000\031\000\000\000\035\000\000\000$\000\000\0002902934.1Everestmountain\004\000\000\000\024\000\000\000\032\000\000\000%\000\000\000)\000\000\000heighthike_safetynametype
[ RECORD 2 ]+-----
__identity__ | 2
__raw__      | \001\000\000\000\000\000\000\004\000\000\000\024\000\000\000\031\000\000\000\035\000\000\000)\000\000\0002902915.4Mt St
Helensvolcano\004\000\000\000\024\000\000\000\032\000\000\000%\000\000\000)\000\000\000heighthike_safetynametype
[ RECORD 3 ]+-----
__identity__ | 3
__raw__      | \001\000\000\000+\000\000\000\004\000\000\000\024\000\000\000\031\000\000\000\035\000\000\000#\000\000\0001700012.2Denalimountain\004\000\000\000\024\000\000\000\032\000\000\000%\000\000\000)\000\000\000heighthike_safetynametype
[ RECORD 4 ]+-----
__identity__ | 4
__raw__      | \001\000\000\000(\000\000\000\003\000\000\000\020\000\000\000\025\000\000\000\000\000\0001400Kilimanjaromountain\003\000\000\000\020\000\000\000\026\000\000\000\032\000\000\000heighthikename
[ RECORD 5 ]+-----
__identity__ | 5
__raw__      | \001\000\000\000)\000\000\000\003\000\000\000\020\000\000\000\024\000\000\000\000\000\00050.6Mt Washingtonmountain\003\000\000\000\020\000\000\000\033\000\000\000\037\000\000\000hike_safetynametype
```

2. Use the `mapToString()` function (with the `__raw__` column of `mountains`) to inspect its contents in readable JSON text format:

```
=> SELECT maptostring(__raw__) from mountains;
MAPTOSTRING
-----
{
  "hike_safety": "50.6",
  "name": "Mt Washington",
  "type": "mountain"
```

```
}
{
  "height":      "29029",
  "hike_safety": "34.1",
  "name": "Everest",
  "type": "mountain"
}
{
  "height":      "14000",
  "hike_safety": "22.8",
  "name": "Kilimanjaro",
  "type": "mountain"
}
{
  "height":      "29029",
  "hike_safety": "15.4",
  "name": "Mt St Helens",
  "type": "volcano"
}
{
  "height":      "17000",
  "hike_safety": "12.2",
  "name": "Denali",
  "type": "mountain"
}
```

3. Now, use the `compute_flextable_keys` function to populate the `mountain_keys` table. Vertica generates this table automatically when you create your flex table.

```
=> SELECT compute_flextable_keys('mountains');
       compute_flextable_keys
-----
Please see public.mountains_keys for updated keys
(1 row)
```

4. Query the keys table `mountains_keys`, and examine the results:

```
=> SELECT * from public.mountains_keys;
 key_name | frequency | data_type_guess
-----+-----+-----
 hike_safety |          5 | varchar(20)
    name    |          5 | varchar(26)
    type    |          5 | varchar(20)
   height   |          4 | varchar(20)
(4 rows)
```

## Build a Flex Table View

1. Use the `build_flextable_view` function to populate a view generated from the `mountains_keys` table.

```
=> SELECT build_flextable_view('mountains');
          build_flextable_view
-----
The view public.mountains_view is ready for querying
(1 row)
```

2. Query the view `mountains_view`:

```
=> SELECT * from public.mountains_view;
 hike_safety |      name      | type  | height
-----+-----+-----+-----
 50.6       | Mt Washington | mountain |
 34.1       | Everest      | mountain | 29029
 22.8       | Kilimanjaro  | mountain | 14000
 15.4       | Mt St Helens | volcano  | 29029
 12.2       | Denali       | mountain | 17000
(5 rows)
```

3. Use the `view_columns` system table to query the `column_name` and `data_type` columns for `mountains_view`:

```
=> SELECT column_name, data_type from view_columns where table_name = 'mountains_view';
 column_name | data_type
-----+-----
 hike_safety | varchar(20)
 name       | varchar(26)
 type       | varchar(20)
 height    | varchar(20)
(4 rows)
```

4. Review the query results:
  - Notice the `data_type` column, its values and sizes. These are calculated when you compute keys for your flex table with `compute_flextable_keys()`.
  - Did you also notice the `data_type_guess` column when you queried the `mountains_keys` table after invoking that function?
5. With the `data_type` information from `mountains_view`, override the `data_type_guess` for `hike_safety`. Then, `COMMIT` the change, and rebuild the view with

`build_flextable_view()`:

```
=> UPDATE mountains_keys SET data_type_guess = 'float' where key_name = 'hike_safety';
OUTPUT
-----
      1
(1 row)

=> commit;
=> SELECT build_flextable_view('mountains');
          build_flextable_view
-----
The view public.mountains_view is ready for querying
(1 row)
```

6. Next, use the `view_columns` system table. Notice that `hike_safety` is now cast to a float data type:

```
=> SELECT column_name, data_type from view_columns where table_name = 'mountains_view';
column_name | data_type
-----+-----
hike_safety | float
name        | varchar(26)
type        | varchar(20)
height      | varchar(20)
(4 rows)
```

## Create a Hybrid Flex Table

If you already know that some of the data you load and query regularly needs full Vertica performance and support, you can create a *hybrid* flex table. A hybrid flex table has one or more real columns that you define, and a `__raw__` column to store any unstructured data you load. Querying real columns is faster than querying flexible data in the `__raw__` column. You can define default values for the columns.

1. Create a hybrid flex table, and load the same sample JSON file:

```
=> CREATE flex table mountains_hybrid(name varchar(41) default name::varchar(41), hike_safety
float
default hike_safety::float);
=> COPY mountains_hybrid from '/home/dbadmin/Downloads/mountains.json' parser fjsonparser();
Rows Loaded
-----
      5
(1 row)
```

2. Use the `compute_flextable_keys_and_build_view` function to populate the

keys table and build the view for mountains\_hybrid:

```
=> SELECT compute_flextable_keys_and_build_view('mountains_hybrid');
                                compute_flextable_keys_and_build_view

-----
--
Please see public.mountains_hybrid_keys for updated keys
The view public.mountains_hybrid_view is ready for querying
(1 row)
```

3. Query the mountains\_hybrid keys table. Review the data\_type\_guesses column values again. The types list the column definitions you declared when you created the hybrid table:

```
=> SELECT * from mountains_hybrid_keys;
key_name | frequency | data_type_guess
-----+-----+-----
height   |          4 | varchar(20)
name     |          5 | varchar(41)
type     |          5 | varchar(20)
hike_safety |          5 | float
(4 rows)
```

If you create a basic flex table, and later find you want to promote one or more virtual columns to real columns, see [Materializing Flex Tables](#) to add columns.

## Promote Virtual Columns in a Hybrid Flex Table

After you explore your flex table data, you can promote one or more virtual columns in your flex table to real columns. You do not need to create a separate columnar table.

1. Invoke the `materialize_flextable_columns` function on the hybrid table, specifying the number of virtual columns to materialize:

```
=> SELECT materialize_flextable_columns('mountains_hybrid', 3);
                                materialize_flextable_columns

-----
The following columns were added to the table public.mountains_hybrid:
type
For more details, run the following query:
SELECT * FROM v_catalog.materialize_flextable_columns_results WHERE table_schema = 'public'
and
```

```
table_name = 'mountains_hybrid';  
  
(1 row)
```

2. You specified three (3) columns to materialize, but the table was created with two real columns (name and hike\_safety). To fulfill your three-column specification, the function promotes only one other column, type. The next example has expanded display to list the columns vertically. Notice the ADDED status for the column that was just materialized, rather than EXISTS for the two columns you defined when creating the table:

```
=> \x  
Expanded display is on.  
=> SELECT * from materialize_flextable_columns_results where table_name = 'mountains_hybrid';  
-[ RECORD 1 ]-----  
table_id      | 45035996273766044  
table_schema  | public  
table_name    | mountains_hybrid  
creation_time | 2013-11-30 20:09:37.765257-05  
key_name      | type  
status        | ADDED  
message       | Added successfully  
-[ RECORD 2 ]-----  
table_id      | 45035996273766044  
table_schema  | public  
table_name    | mountains_hybrid  
creation_time | 2013-11-30 20:09:37.765284-05  
key_name      | hike_safety  
status        | EXISTS  
message       | Column of same name already exists in table definition  
-[ RECORD 3 ]-----  
table_id      | 45035996273766044  
table_schema  | public  
table_name    | mountains_hybrid  
creation_time | 2013-11-30 20:09:37.765296-05  
key_name      | name  
status        | EXISTS  
message       | Column of same name already exists in table definition
```

3. Now, display the hybrid table definition, listing the \_\_raw\_\_ column and the three materialized columns. Flex table data types are derived from the associated keys tables, so you can update them as necessary. Notice that the \_\_raw\_\_ column has a default NOT NULL constraint:

```
=> \d mountains_hybrid  
List of Fields by Tables  
-[ RECORD 1 ]-----  
Schema      | public  
Table       | mountains_hybrid  
Column      | __raw__  
Type        | long varbinary(130000)
```

## Using Flex Tables

### Getting Started

```
Size          | 130000
Default       |
Not Null      | t
Primary Key   | f
Foreign Key   |
-[ RECORD 2 ]-----
Schema        | public
Table         | mountains_hybrid
Column        | name
Type          | varchar(41)
Size          | 41
Default       | (MapLookup(mountains_hybrid.__raw__, 'name'))::varchar(41)
Not Null      | f
Primary Key   | f
Foreign Key   |
-[ RECORD 3 ]-----
Schema        | public
Table         | mountains_hybrid
Column        | hike_safety
Type          | float
Size          | 8
Default       | (MapLookup(mountains_hybrid.__raw__, 'hike_safety'))::float
Not Null      | f
Primary Key   | f
Foreign Key   |
-[ RECORD 4 ]-----
Schema        | public
Table         | mountains_hybrid
Column        | type
Type          | varchar(20)
Size          | 20
Default       | (MapLookup(mountains_hybrid.__raw__, 'type'))::varchar(20)
Not Null      | f
Primary Key   | f
Foreign Key   |
```

You have now completed getting started with flex table basics, hybrid flex tables, and using the flex functions.





# Understanding Flex Tables

You can create flex tables and then manage them with their associated helper, data, and map functions. Flex tables:

- Do not require schema definitions
- Do not need column definitions
- Have full Unicode support
- Support SQL queries

You can use flex tables to promote data directly from exploration to analytic operations. flex tables features include:

- Ability to load different formats into one flex table, which lets you handle changing structure over time
- Full support of delimited and JSON data
- Extensive SQL queries and built-in analytics for the data you load
- Usability functions, which let you explore your unstructured data and then use built-in functions to materialize the data

## Exploration to Promotion

After you create a flex table, you can quickly load data, including social media content in JSON, log files, delimited files, and other information. Previously, working with such data required significant schema design and preparation. Now, you can load and query flex tables in a few steps.

Creating flex tables is similar to creating other tables, except column definitions are optional. When you create flex tables, with or without column definitions, Vertica implicitly adds a real column to your table, called `__raw__`. This column stores loaded data. The `__raw__` column is a `LONG VARBINARY` column with a `NOT NULL` constraint. It contains the documented limits for its data type (see [Long Data Types](#) in the SQL Reference Manual). The `__raw__` column's default maximum width is 130,000 bytes (with an absolute maximum of 32,000,000 bytes). You can change the default width with the `FlexTablesRawSize` configuration parameter.

If you create a flex table without other column definitions, the table includes a second default column, `__identity__`, declared as an auto-incrementing `IDENTITY (1,1)`

column. When no other columns are defined, flex tables use the `__identity__` column for segmentation and sort order.

Loading data into a flex table encodes the record into a VMap type and populates the `__raw__` column. The VMap is a standard dictionary type, pairing keys with string values as virtual columns.

## Flex Table Terms

This guide uses the following terms when describing how you work with flex tables to explore and analyze flexible data:

- **VMap**: An internal map data format.
- **Virtual Columns**: Undeclared columns that have not been promoted to real columns.
- **Real Columns**: Fully featured columns in flex or columnar tables.
- **Promoted Columns**: Virtual columns that have been materialized to real columns.
- **Map Keys**: Map keys are the virtual column names within VMap data.

## Is There Structure in a Flex Table?

The term *unstructured data* (sometimes called *semi-structured* or *Dark Data*) does not indicate that the data you load into flex tables is entirely without structure. However, you may not know the data's composition or the inconsistencies of its design. In some cases, the data may not be relational.

Your data may have some structure (like JSON and delimited data). Data may be semi-structured or stringently structured, but in ways that you either do not know about or do not expect. In this guide, the term *flexible data* encompasses these and other kinds of data. You can load your flexible data directly into a flex table, and query its contents with your favorite SQL SELECT or other statements.

To summarize, you can load data first, without knowing its structure, and then query its content after a few simple transformations. In some cases, you already know the data structure, such as some tweet map keys, like `user.lang`, `user.screen_name`, and `user.url`. If so, you can query these values explicitly as soon as you load the data.

## Making Flex Table Data Persist

The underlying implementation of each flex table is one (or two) real columns. Because of this design, existing Vertica functionality writes the table and its contents to disk

(ROS). This approach maintains K-safety in your cluster and supports standard recovery processes should node failures occur. Flex tables are included in full backups (or, if you choose, in object-level backups).

## What Happens When You Create Flex Tables?

Whenever you execute a `CREATE FLEX TABLE` statement, Vertica creates three objects, as follows:

- The flexible table (*flex\_table*)
- An associated keys table (*flex\_table\_keys*)
- A default view for the main table (*flex\_table\_view*)

The `_keys` and `_view` objects are dependents of the parent, *flex\_table*. Dropping the flex table also removes its dependents, although you can drop the `_keys` or `_view` objects independently.

You can create a flex table without specifying any column definitions (such as `darkdata`, in the next example). When you do so, Vertica automatically creates two tables, the named flex table (such as `darkdata`) and its associated keys table, `darkdata_keys`:

```
=> create flex table darkdata();
CREATE TABLE
=> \dt dark*

      List of tables
Schema |      Name      | Kind | Owner  | Comment
-----+-----+-----+-----+-----
public | darkdata       | table | dbadmin |
public | darkdata_keys  | table | dbadmin |
(2 rows)
```

Each flex table has two default columns, `__raw__` and `__identity__`. The `__raw__` column exists in every flex table to hold the data you load. The `__identity__` column is auto-incrementing. Vertica uses the `__identity__` column for segmentation and sort order when no other column definitions exist. The flex keys table (`darkdata_keys`) has three columns, as shown:

```
=> select * from darkdata;
__identity__ | __raw__
-----+-----
(0 rows)

=> select * from darkdata_keys;
```

## Using Flex Tables

### Understanding Flex Tables

```
key_name | frequency | data_type_guess
-----+-----+-----
(0 rows)
```

Creating a flex table with column definitions (such as `darkdata1`, in the next example) automatically generates a table with the `__raw__` column. However, the table has no `__identity__` column because columns are specified for segmentation and sort order. Two tables are created automatically, as shown in the following example:

```
=> create flex table darkdata1 (name varchar);
CREATE TABLE

=> select * from darkdata1;
__raw__ | name
-----+-----
(0 rows)

=> \d darkdata1*

                                List of Fields by Tables
 Schema | Table | Column | Type | Size | Default | Not Null | Primary Key
 | Foreign Key
-----+-----+-----+-----+-----+-----+-----+-----
+-----+
public | darkdata1 | __raw__ | long varbinary(130000) | 130000 | | t | f
|
public | darkdata1 | name | varchar(80) | 80 | | f | f
|
(2 rows)

=> \dt darkdata1*

                                List of tables
 Schema | Name | Kind | Owner | Comment
-----+-----+-----+-----+-----
public | darkdata1 | table | dbadmin |
public | darkdata1_keys | table | dbadmin |
(2 rows)
```

Creating a flex table with at least one column definition (`darkdata1` in the next example) also generates a table with the `__raw__` column, but not an `__identity__` column. Instead, the specified columns are used for segmentation and sort order. Two tables are also created automatically, as shown in the following example:

```
=> create flex table darkdata1 (name varchar);
CREATE TABLE

=> \d darkdata1*

                                List of Fields by Tables
 Schema | Table | Column | Type | Size | Default | Not Null | Primary Key
 | Foreign Key
-----+-----+-----+-----+-----+-----+-----+-----
+-----+
public | darkdata1 | __raw__ | long varbinary(130000) | 130000 | | t | f
|
public | darkdata1 | name | varchar(80) | 80 | | f | f
|
(2 rows)
```

```
=> \dt darkdata1*

      List of tables
 Schema |      Name      | Kind | Owner  | Comment
-----+-----+-----+-----+-----
 public | darkdata1       | table | dbadmin |
 public | darkdata1_keys | table | dbadmin |
(2 rows)
```

For more examples, see [Creating Flex Tables](#).

## Creating Superprojections Automatically

In addition to creating two tables for each flex table, Vertica creates superprojections for both the main flex table and its associated keys table. Using the `\dj` command, you can display the projections created automatically for the `darkdata` and `darkdata1` tables in this set of examples:

```
=> \dj darkdata*

      List of projections
 Schema |      Name      | Owner  | Node           | Comment
-----+-----+-----+-----+-----
 public | darkdata1_b0   | dbadmin |                |
 public | darkdata1_b1   | dbadmin |                |
 public | darkdata1_keys_super | dbadmin | v_vmart_node0001 |
 public | darkdata1_keys_super | dbadmin | v_vmart_node0003 |
 public | darkdata1_keys_super | dbadmin | v_vmart_node0004 |
 public | darkdata_b0    | dbadmin |                |
 public | darkdata_b1    | dbadmin |                |
 public | darkdata_keys__super | dbadmin | v_vmart_node0001 |
 public | darkdata_keys_super | dbadmin | v_vmart_node0003 |
 public | darkdata_keys_super | dbadmin | v_vmart_node0004 |
(10 rows)
```

**Note:** You cannot create pre-join projections from flex tables. For information about projections and how they are used, see [Vertica Concepts](#).

## Default Flex Table View

When you create a flex table, you also create a default view. This default view uses the table name with a `_view` suffix, as listed in the next example, which shows the list of views for `darkdata` and `darkdata1`. If you query the default view, you are prompted to use the [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#) function. This view enables you to update the view after you load data so that it includes all keys and values.

```
=> \dv darkdata*

      List of View Fields
 Schema |      View      | Column | Type  | Size
-----+-----+-----+-----+-----
```

```
public | darkdata_view | status | varchar(124) | 124  
public | darkdata1_view | status | varchar(124) | 124  
(2 rows)
```

For more information, see [Updating Flex Table Views](#).

## Flex Functions

There are three sets of functions to support flex tables and extracting data into VMaps. See the following sections for more information:

- Data (helper) functions ([Flex Data Functions Reference](#))
- Extractor functions ([Flex Extractor Functions Reference](#))
- Map functions ([Flex Map Functions Reference](#))

## Using Clients with Flex Tables

You can use the Vertica supported client drivers with flex tables as follows:

- Flex tables do not permit `INSERT` statements from `vsq`l or from any client. To load data from a client, use `COPY LOCAL` with the appropriate flex table parser.
- The driver metadata APIs return only real columns. For example, using a `select * from myflex;` statement, when *myflex* has a single materialized column (`name`), returns the `__raw__` and `name` columns. However, it does not return virtual columns from within `__raw__`. To access virtual columns and their values, query the associated `flextable_keys` table, just as you would in `vsq`l.

## Creating Flex Tables

You can create a flex table, or an external flex table, without column definitions or other parameters. You can use any CREATE TABLE statement parameters you prefer, as usual.

## Unsupported CREATE FLEX TABLE Statements

These statements are not currently supported:

- CREATE FLEX TABLE AS...
- CREATE FLEX TABLE LIKE...

## Creating Basic Flex Tables

Here's how to create the table:

```
=> create flex table darkdata();  
CREATE TABLE
```

Selecting from the table before loading any data into it reveals its two real columns, `__identity__` and `__raw__`:

```
=> select * from darkdata;  
__identity__ | __raw__  
-----+-----  
(0 rows)
```

Here's an example of creating a flex table with a column definition:

```
=> create flex table darkdata1(name varchar);  
CREATE TABLE
```

When flex tables exist, you can add new columns (including those with default derived expressions), as described in [Altering Flex Tables](#).

## Creating Temporary Flex Tables

Before you create temporary global and local flex tables, be aware of the following considerations:

- GLOBAL TEMP flex tables are supported. Creating a temporary global flex table results in the `flextable_keys` table and the `flextable_view` having temporary table restrictions for their content.
- LOCAL TEMP flex tables must include at least one column definition. The reason for this requirement is that local temp tables do not support automatically-incrementing data (such as the flex table default `__identity__` column). Creating a temporary local flex table results in the `flextable_keys` table and the `flextable_view` existing in the local temporary object scope.
- LOCAL TEMP views are supported for flex and columnar temporary tables.

For global or local temp flex tables to function correctly, you must also specify the `ON COMMIT PRESERVE ROWS` clause. You must use the `ON COMMIT` clause for the flex table helper functions, which rely on commits. Create a local temp table as follows:

```
=> create flex local temp table good(x int) ON COMMIT PRESERVE ROWS;  
CREATE TABLE
```

After creating a local temporary flex table using this approach, you can then load data into the table, create table keys, and a flex table view:

```
=> copy good from '/home/release/KData/bake.json' parser fjsonparser();  
Rows Loaded  
-----  
1  
(1 row)  
  
=> select compute_flextable_keys_and_build_view('good');  
compute_flextable_keys_and_build_view  
-----  
Please see v_temp_schema.good_keys for updated keys  
The view good_view is ready for querying  
(1 row)
```

However, creating temporary flex tables without an `ON COMMIT PRESERVE ROWS` clause results in the following warnings:

```
=> create flex local temp table bak1(id int, type varchar(10000), name varchar(1000));  
WARNING 5860: Due to the data isolation of temp tables with an on-commit-delete-rows  
policy, the compute_flextable_keys() and compute_flextable_keys_and_build_view() functions  
cannot access this table's data. The build_flextable_view() function can be used with a  
user-provided keys table to create a view, but involves a DDL commit which will delete  
the table's rows
```

After loading data into a such a temporary flex table, computing keys or building a view for the flex table results in the following error:

```
=> select compute_flextable_keys('bak1');  
ERROR 5859: Due to the data isolation of temp tables with an on-commit-delete-rows policy,
```



```
the compute_flextable_keys() and compute_flextable_keys_and_build_view() functions cannot
access this table's data
HINT: Make the temp table ON COMMIT PRESERVE ROWS to use this function
```

Similarly, you can create global temp tables as follows:

```
=> create flex global temp table good_global(x int) ON COMMIT PRESERVE ROWS;
```

After creating a global temporary flex table using this approach, you can then load data into the table, create table keys, and a flex table view:

```
=> copy good_global from '/home/dbadmin/data/flex/bake_single.json' parser fjsonparser();
Rows Loaded
-----
5
(1 row)

=> select compute_flextable_keys_and_build_view('good_global');
           compute_flextable_keys_and_build_view
-----
Please see v_temp_schema.good_keys for updated keys
The view good_view is ready for querying
(1 row)
```

Similar to local temp flex tables, creating global flex tables without an `ON COMMIT PRESERVE ROWS` clause results in the following warnings:

```
=> create flex global temp table bak_global(id int, type varchar(10000), name varchar(1000));
WARNING 5860: Due to the data isolation of temp tables with an on-commit-delete-rows
policy, the compute_flextable_keys() and compute_flextable_keys_and_build_view() functions
cannot access this table's data. The build_flextable_view() function can be used with a
user-provided keys table to create a view, but involves a DDL commit which will delete
the table's rows
CREATE TABLE
```

Loading data into a such a temporary flex table, computing keys or building a view for the flex table results in the following error:

```
=> select compute_flextable_keys('bak_global');
ERROR 5859: Due to the data isolation of temp tables with an on-commit-delete-rows policy, the
compute_flextable_keys() and
compute_flextable_keys_and_build_view() functions cannot access this table's data
HINT: Make the temp table ON COMMIT PRESERVE ROWS to use this function
```

## Materializing Flex Table Virtual Columns

After you create your flex table and load data, you compute keys from virtual columns. After completing those tasks, you can materialize some keys by promoting virtual columns to real table columns. By promoting virtual columns, you query real columns rather than the raw data.

You can promote one or more virtual columns — materializing those keys from within the `__raw__` data to real columns. Vertica recommends this approach to get the best query performance for all important keys. You don't need to create new columnar tables from your flex table.

Promoting flex table columns results in a hybrid table. Hybrid tables:

- Maintain the convenience of a flex table for loading unstructured data
- Improve query performance for any real columns

If you have only a few columns to materialize, try altering your flex table progressively, adding columns whenever necessary. You can use the `ALTER TABLE...ADD COLUMN` statement to do so, just as you would with a columnar table. See [Altering Flex Tables](#) for ideas about adding columns.

If you want to materialize columns automatically, use the helper function [MATERIALIZED\\_FLEXTABLE\\_COLUMNS](#)

## Creating Columnar Tables from Flex Tables

You can create a regular columnar Vertica table from a flex table, but you cannot use one flex table to create another.

Typically, you create a columnar table from a flex table after loading data. Then, you specify the virtual column data you want in a regular table, casting virtual columns to regular data types.

To create a columnar table from a flex table, `darkdata`, select two virtual columns, (`user.lang` and `user.name`), for the new table. Use a command such as the following, which casts both columns to `varchar`s for the new table:

```
=> create table darkdata_full as select "user.lang"::varchar, "user.name"::varchar from darkdata;
CREATE TABLE
=> select * from darkdata_full;
 user.lang |      user.name
-----+-----
 en        | Frederick Danjou
 en        | The End
 en        | Uptown gentleman.
 en        | ~G A B R I E L A â
 es        | Flu Beach
 es        | I'm Toasterâ
 it        | laughing at clouds.
 tr        | seydo shi
          |
          |
          |
(12 rows)
```

## Creating External Flex Tables

To create an external flex table:

```
=> create flex external table mountains() as copy from 'home/release/KData/kmm_ountains.json'
parser fjsonparser();
CREATE TABLE
```

As with other flex tables, creating an external flex table produces two regular tables: the named table and its associated `_keys` table. The keys table is not an external table:

```
=> \dt mountains
List of tables
Schema | Name | Kind | Owner | Comment
-----+-----+-----+-----+-----
public | mountains | table | release |
(1 row)
```

You can use the helper function, [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#), to compute keys and create a view for the external table:

```
=> select compute_flextable_keys_and_build_view ('appLog');
```

```
compute_flextable_keys_and_build_view
-----
Please see public.appLog_keys for updated keys
The view public.appLog_view is ready for querying
(1 row)
```

1. Check the keys from the `_keys` table for the results of running the helper application:

```
=> select * from appLog_keys;
key_name | frequency | data_type_guess
-----+-----+-----
contributors | 8 | varchar(20)
coordinates | 8 | varchar(20)
created_at | 8 | varchar(60)
entities.hashtags | 8 | long varbinary(186)
.
.
.
retweeted_status.user.time_zone | 1 | varchar(20)
retweeted_status.user.url | 1 | varchar(68)
retweeted_status.user.utc_offset | 1 | varchar(20)
retweeted_status.user.verified | 1 | varchar(20)
(125 rows)
```

2. Query from the external flex table view:

```
=> select "user.lang" from appLog_view;
```

```
user.lang
-----
it
en
es
en
en
es
tr
en
(12 rows)
```

**Note:** External tables are fully supported for both flex and columnar tables. However, using external flex (or columnar) tables is less efficient than using flex tables whose data is stored in the Vertica database. Data that is maintained externally requires reloading each time you query..

## Partitioning Flex Tables

You cannot partition a flex table on any virtual column (key).

The next example shows a query on `user.location`, which is a virtual column in the map data. The example then attempts to partition that column:

```
=> select "user.location" from darkdata;
user.location
-----
chicago
Narnia
Uptown..
Chile
(12 rows)
=> alter table darkdata partition by "user.location" reorganize;
ROLLBACK 5371: User defined function not allowed: MapLookup
```

# Using COPY with Flex Tables

You load data into a flex table with a COPY statement, specifying one of the flex parsers:

- FAVROPARSER
- FCEFPARSER
- FCSVPARSER
- FDELIMITEDPAIRPARSER
- FDELIMITEDPARSER
- FJSONPARSER
- FREGEXPARSER

All flex parsers store the data as a single-value VMap. They reside in the VARBINARY `__raw__` column, which is a real column with a NOT NULL constraint. The VMap is encoded into a single binary value for storage in the `__raw__` column. The encoding places the value strings in a contiguous block, followed by the key strings. Vertica supports null values within the VMap for keys with NULL-specified columns. The key and value strings represent the virtual columns and their values in your flex table.

If a flex table data row is too large to fit in the VARBINARY `__raw__` column, it is rejected. By default, the rejected data and exceptions files are stored in the standard CopyErrorLogs location, a subdirectory of the catalog directory:

```
v_mart_node003_catalog\CopyErrorLogs\trans-STDIN-copy-from-exceptions.1  
v_mart_node003_catalog\CopyErrorLogs\trans-STDIN-copy-rejections.1
```

Flex tables do not copy any rejected data, due to disk space considerations. The rejected data file exists, but it contains only a new line character for every rejected record. The corresponding exceptions file lists the reason why each record was rejected.

You can specify a different path and file for the rejected data and exceptions files. To do so, use the COPY parameters REJECTED DATA and EXCEPTIONS, respectively. You can also save load rejections and exceptions in a table. See the *Bulk Loading* section of the Administrator's Guide, (*Capturing Load Rejections and Exceptions*).

## Basic Flex Table Load and Query

Loading data into your flex table is similar to loading data into a regular columnar table. The difference is that you must use the parser argument with one of the flex parsers:

```
=> copy darkdata from '/home/dbadmin/data/tweets_12.json' parser fjsonparser();  
Rows Loaded  
-----  
12  
(1 row)
```

**Note:** You can use many additional COPY parameters as required but not all are supported.

## Loading Data into Flex Table Real Columns

If you create a hybrid flex table with one or more real column definitions, COPY evaluates each virtual column key name during data load. For each real column with a name that is identical to a virtual column key name, COPY does the following:

- Loads the keys and values as part of the VMap data in the `__raw__` column
- Automatically populates real columns with the values from their virtual column counterparts

Subsequent data loads continue loading same-name key-value pairs into both the `__raw__` column and the real column.

**Note:** Over time, storing values in both column types can impact your licensed data limits. For more information about Vertica licenses, see [Managing Licenses](#) in the Administrator's Guide.

For example, continuing with the JSON data:

1. Create a flex table, `darkdata1`, with a column definition of one of the keys in the data you will load:

```
=> create flex table darkdata1 ("user.lang" varchar);  
CREATE TABLE
```

2. Load data into `darkdata1`:

```
=> copy darkdata1 from '/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
                12
(1 row)
```

3. Query the `user.lang` column of `darkdata1`. Loading the JSON data file populated the column you defined:

```
=> select "user.lang" from darkdata1;
user.lang
-----
es
es
tr
it
en
en
en
en
en
(12 rows)
```

Empty column rows indicate NULL values. For more information about how NULLs are handled in flex table, see [\(\)](#).

4. You can query for other virtual columns (such as `"user.name"` in `darkdata1`), with similar results as for `"user.lang"`:

```
=> select "user.name" from darkdata1;
user.name
-----
I'm Toasterâ
Flu Beach
seydo shi
The End
Uptown gentleman.
~G A B R I E L A â
Frederick Danjou
laughing at clouds.
(12 rows)
```

**Note:** While the results for these two queries are similar, the difference in accessing the keys and their values is significant. Data for `"user.lang"` has been materialized into a real table column, while `"user.name"` remains a virtual column. For production-level data usage (rather than test data sets), materializing flex table data improves query performance significantly.

## Handling Default Values During Loading

You can create your flex table with a real column, named for a virtual column that exists in your incoming data. For example, if the data you load has a `user.lang` virtual column, define the flex table with that column. You can also specify a default column value when creating the flex table with real columns. The next example shows how to define a real column (`user.lang`), which has a default value from a virtual column (`user.name`):

```
=> create flex table table darkdata1 ("user.lang" long varchar default "user.name");
```

When you load data into your flex table, COPY uses values from the flex table data, ignoring the default column definition. Loading data into a flex table requires **MAPLOOKUP** to find keys that match any real column names. A match exists when the incoming data has a virtual column with the same name as a real column. When COPY detects a match, it populates the column with values. COPY returns either a value or NULL for each row, so real columns always have values.

For example, after creating the `darkdata1` flex table, described in the previous example, load data with COPY:

```
=> copy darkdata1 from '/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
          12
(1 row)
```

If you query the `darkdata1` table after loading, the data shows that the values for the `user.lang` column were extracted:

- From the data being loaded — Values for the `user.lang` virtual column
- With NULL — Rows without values

In this case, the table column default value for `user.lang` was ignored:

```
=> select "user.lang" from darkdata1;
user.lang
-----
it
en
es
en
en
es
tr
en
(12 rows)
```



## Using COPY to Specify Default Column Values

You can add an expression to a COPY statement to specify default column values when loading data. For flex tables, specifying any column information requires that you list the `__raw__` column explicitly. The following example shows how to use an expression for the default column value. In this case, loading populates the defined `user.lang` column with data from the input data `user.name` values:

```
=> copy darkdata1(__raw__, "user.lang" as "user.name"::varchar)from
'/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
          12
(1 row)
=> select "user.lang" from darkdata1;
      user.lang
-----
laughing at clouds.
Avita Desai
I'm Toasterâ
Uptown gentleman.
~G A B R I E L A â
Flu Beach
seydo shi
The End
(12 rows)
```

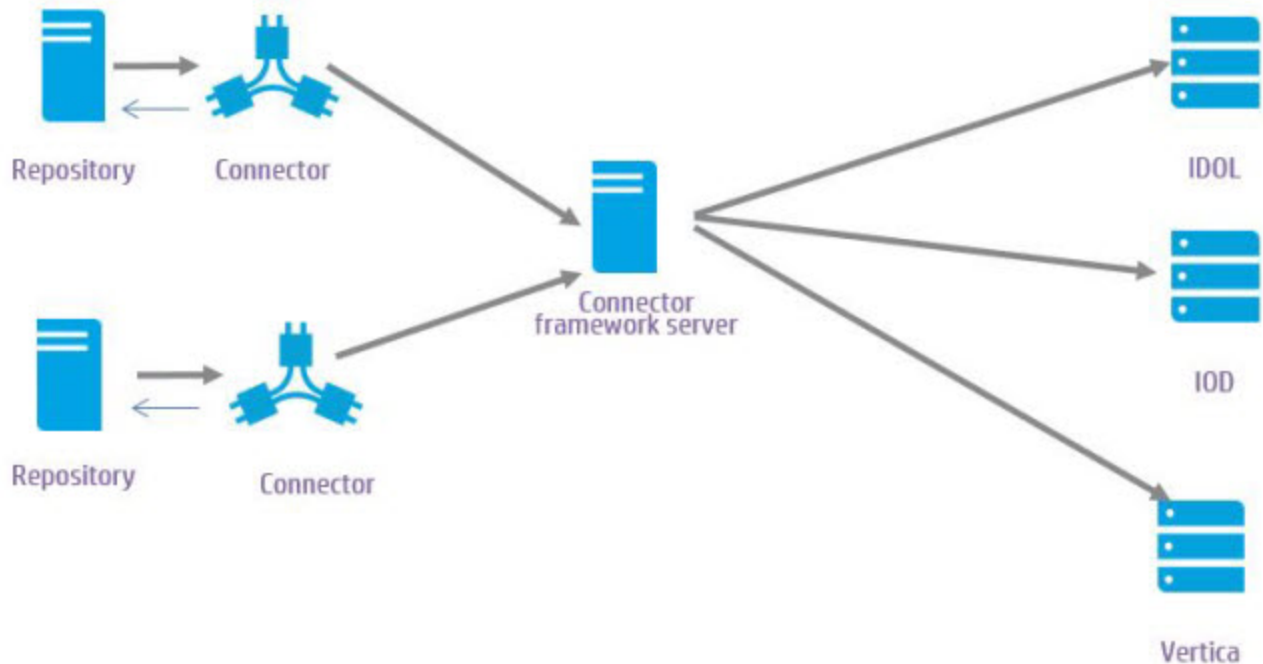
You can specify default values when adding columns, as described in [Altering Flex Tables](#). When you do so, a different behavior results. For more information about using COPY, its expressions and parameters, see [Bulk Loading Data](#) in the Administrator's Guide and [COPY](#) in the SQL Reference Manual.

## Using Flex Tables for IDOL Data

You can create flex tables to use with the HPE IDOL Connector Framework Server (CFS) and an ODBC client. The CFS VerticalIndexer module uses the connector to retrieve data. CFS then indexes the data into your Vertica database.

CFS supports many connectors for interfacing to different unstructured file types stored in repositories. Examples of repositories include Microsoft Exchange (email), file systems (including Word documents, images, and videos), Microsoft SharePoint, and Twitter (containing Tweets).

Connectors retrieve and aggregate data from repositories. CFS indexes the data, sending it to IDOL, IDOL OnDemand, or Vertica. The following figure illustrates a basic setup with a repository and a connector.



After you configure CFS and connect it to your Vertica database, the connector monitors the repository for changes and deletions to loaded documents, and for new files not previously added to the server. CFS then updates its server destinations automatically.

To achieve the best query results with ongoing CFS updates and deletes, HPE recommends using live aggregate projections and top-K projections. For more information about how these projections work, and for examples of using them, see [Working with Projections](#) in the Administrator's Guide.

## ODBC Connection String for CFS

There are several steps to setting up the CFS VerticalIndexer to load IDOL metadata into your database. For a more complete example, see the blog post [HPE IDOL CFS Vertica Module](#).

One of the first steps is to add information to the CFS configuration file. To do so, add an Indexing section to the configuration file that specifies the ODBC ConnectionString details.

Successfully loading data requires a valid database user with write permissions to the destination table. Two ODBC connection parameters (UID and PWD) specify the Vertica user and password. The following example shows a sample CFS Indexing section. The section includes a ConnectionString with the basic parameters, including a sample user (UID=fjones) and password (PWD=fjones\_password):

```
[Indexing]
IndexerSections=vertica
IndexTimeInterval=30

[vertica]
IndexerType = Library
ConnectionString=Driver=Vertica;Server=123.456.478.900;Database=myDB;UID=fjones;PWD=fjones_password
TableName = marcomm.myFlexTable
LibraryDirectory = ./shared_library_indexers
LibraryName = verticaIndexer
```

For more information about ODBC connection parameters, see [ODBC Configuration Parameters](#).

## CFS COPY LOCAL Statement

CFS first indexes and processes metadata from a document repository to add to your database. Then, CFS uses the Indexing information you added to the configuration file to create an ODBC connection. After establishing a connection, CFS generates a standard COPY LOCAL statement, specifying the fjsonparser. CFS loads data directly into your pre-existing flex table with a statement such as the following:

```
=> COPY myFlexTable FROM LOCAL path_to_compressed_temporary_json_file parser fjsonparser();
```

```
=> select * from myavro;
__identity__ | __raw__
-----+-----
(0 rows)
```

When your IDOL metadata appears in a flex table, you can optionally add new table columns, or materialize other data, as described in [Altering Flex Tables](#).

# Using Flex Table Parsers

You can load flex tables with one of several parsers. You can load data using the options that the flex parsers support:

- [Loading Avro Data](#)
- [Loading Common Event Format \(CEF\) Data](#)
- [Loading CSV Data](#)
- [Loading Delimited Data](#)
- [Loading JSON Data](#)
- [Loading Matches from Regular Expressions](#)

## Using Flex Parsers for Columnar Tables

You can use any of the flex parsers to load data into columnar tables. Using the flex table parsers to load columnar tables gives you the capability to mix data loads in one table. For example, you can load JSON data in one session and delimited data in another.

**Note:** For Avro data, you can load only data into a columnar table, not the schema. For flex tables, Avro schema information is required to be embedded in the data.

The following basic examples illustrate how you can use flex parsers with columnar tables.

1. Create a columnar table, `super`, with two columns, `age` and `name`:

```
=> create table super(age int, name varchar);  
CREATE TABLE
```

2. Enter JSON values from STDIN, using the `fjsonparser()`.

```
=> copy super from stdin parser fjsonparser();  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> {"age": 5, "name": "Tim"}  
>> {"age": 3}  
>> {"name": "Fred"}  
>> {"name": "Bob", "age": 10}  
>> \.
```

3. Query the table to see the values you entered:

```
=> select * from super;
age | name
-----+-----
    | Fred
  10 | Bob
   5 | Tim
   3 | 
(4 rows)
```

4. Enter some delimited values from STDIN, using the `fdelimitedparser()`:

```
=> copy super from stdin parser fdelimitedparser();
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> name |age
>> Tim|50
>> |30
>> Fred|
>> Bob|100
>> \.
```

5. Query the flex table. Both JSON and delimited data are saved in the same columnar table, `super`.

```
=> select * from super;
age | name
-----+-----
  50 | Tim
  30 | 
   3 | 
   5 | Tim
 100 | Bob
    | Fred
  10 | Bob
    | Fred
(8 rows)
```

Use the `reject_on_materialized_type_error` parameter to avoid loading data with type mismatch. If `reject_on_materialized_type_error` is set to `false`, the flex parser will accept the data with type mismatch. Consider the following example:

Assume that the CSV file to be loaded has the following sample contents:

```
$ cat json.dat
{"created_by":"system","site_source":"flipkart_india_kol","updated_by":"system1","invoice_id":"INVDPKOL100","vendor_id":"VEN15731","total_quantity":12,"created_at":"2012-01-09 23:15:52.0"}
{"created_by":"system","site_source":"flipkart_india_kol","updated_by":"system2","invoice_id":"INVDPKOL101","vendor_id":"VEN15732","total_quantity":14,"created_at":"hello"}
```

1. Create a columnar table.

```
=> CREATE TABLE hdfs_test (  
  site_source varchar(200),  
  total_quantity int ,  
  vendor_id varchar(200),  
  invoice_id varchar(200),  
  updated_by varchar(200),  
  created_by varchar(200),  
  created_at timestamp  
);
```

2. Load JSON data.

```
=> COPY hdfs_test from '/home/dbadmin/json.dat' parser fjsonparser() ABORT ON ERROR;  
Rows Loaded  
-----  
2  
(1 row)
```

3. View the contents.

```
=> select * from hdfs_test;  
site_source | total_quantity | vendor_id | invoice_id | updated_by | created_by | created_at  
-----+-----+-----+-----+-----+-----+-----  
flipkart_india_kol | 12 | VEN15731 | INVDPKOL100 | system1 | system | 2012-01-09 23:15:52  
flipkart_india_kol | 14 | VEN15732 | INVDPKOL101 | system2 | system |  
(2 rows)
```

4. If reject\_on\_materialized\_type\_error parameter is set to true, you will receive errors when loading the sample JSON data.

```
=> COPY hdfs_test from '/home/dbadmin/data/flex/json.dat' parser fjsonparser(reject_on_  
materialized_type_error=true) ABORT ON ERROR;  
ERROR 2035: COPY: Input record 2 has been rejected (Rejected by user-defined parser)
```

## Loading Avro Data

You can load Avro data files into flex tables and columnar tables using the parser, `favroparser`. Before loading, verify that Avro files are encoded in the Avro binary serialization encoding format, described in the [Apache Avro standard](#). The parser also supports Snappy compression. You cannot load Avro data directly from STDIN.

**Note:** The parser `favroparser` does not support Avro files with separate schema files. The Avro file must have its related schema in the file you are loading.

You can use the following data types and optional parameters for `favroparser`.

The `favroparser` supports two data types:

- [Primitive Data Types for favroparser](#)
- [Complex Data Types for favroparser](#)

## Rejecting Data on Materialized Column Type Errors

The favroparser has a Boolean parameter, `reject_on_materialized_type_error`. If you set this parameter to `true`, Vertica rejects rows when the input data presents *both* of the following conditions:

- Includes keys matching an existing materialized column
- Has a value that cannot be coerced into the materialized column's data type

Suppose the flex table has a materialized column, `Temperature`, declared as a `FLOAT`. If you try to load a row with a `Temperature` key that has a `VARCHAR` value, favroparser rejects the data row.

## See Also

- [Using COPY with Kafka](#)

## Primitive Data Types for favroparser

The favroparser supports the following primitive data types:

AVRO Data Type	Vertica Data Type	Value
NULL	<a href="#">NULL</a>	No value
boolean	<a href="#">BOOLEAN</a>	A binary value
int	<a href="#">INTEGER</a>	32-bit signed integer
long	<a href="#">INTEGER</a>	64-bit signed integer
float	<a href="#">DOUBLE PRECISION (FLOAT)</a> Synonymous with 64-bit IEEE FLOAT	Single precision

		(32-bit) IEEE 754 floating- point number
double	DOUBLE PRECISION (FLOAT)	Double precision (64-bit) IEEE 754 floating- point number
bytes	BYTES	Sequence of 8-bit unsigned bytes
string	VARCHAR	Unicode character sequence

**Note:** Vertica does not have an explicit 4-byte (32-bit integer) or smaller types. Instead, Vertica encoding and compression automatically eliminate the storage overhead of values that require less than 64 bits.

Vertica copies each primitive type into the `__raw__` column of the flex table. In this copy operation, the name of the primitive type becomes a virtual column key with its corresponding value as the value of the virtual column.

If the flex table has materialized columns, `favroparser` loads the primitive data type into the corresponding Vertica type for the column. If the parsing is successful, Vertica copies the data into the materialized column; otherwise, it rejects the row.

## Complex Data Types for `favroparser`

You specify the data type of a record in the Avro file using the `type` parameter for `favroparser`. The `favroparser` supports these complex data types:

- [Records](#)
- [Enums](#)
- [Arrays](#)



- [Maps](#)
- [Unions](#)
- [Fixed](#)

This section describes attributes associated with the complex data types.

## Records

Records have the following attributes:

Attribute	Description
name	A JSON string for the name of the record
fields	<p>A JSON array used to list fields. Each field is a JSON object:</p> <ul style="list-style-type: none"><li>• name: A JSON string for the name of the field</li><li>• type: A JSON object used to define a schema or a JSON string used for naming a record definition</li></ul>

The name of each field is used as a virtual column name. If `flatten_records = true` and several nesting levels are present, Vertica concatenates the record names to create the `key_name`, as follows:

```
{
  "type": "record",
  "name": "Profile",
  "fields" : [
    {"name": "UserName", "type": "string"},
    {"name": "Address", "type": "string"}
  ]
}
```

```
{
  "type": "record",
  "name": "Profile",
  "fields" : [
    {VerticaUser},
    {VerticaUser Address}
  ]
}
```

Vertica creates virtual columns for the records as follows:

Names	Values
UserName	VerticaUser
Address	VerticaUser Address

## Enums

Enums (enumerated values) use the type name `enum` and support the following attributes:

Attribute	Description
name	A JSON string for the name of the enum
symbols	A JSON array used to list symbols as JSON strings. All symbols in an enum must be unique and duplicates are prohibited

Example:

```
{
  "type": "enum",
  "name": "suit",
  "symbols" : ["SPADES", "HEARTS", "DIAMONDS", "CLUBS"]
}
```

Consider the preceding Avro schema with a record that contains a field with the value `HEARTS`. In this case, the key value pair copied into the `__raw__` column has `suit` as the key and `HEARTS` as the value.

## Arrays

Arrays use the type name `array` and support one attribute:

Attribute	Description
items	The schema of the array's items

For example, declare an array of strings:

```
{"type": "array", "items": "string"}
```

Similar to the capabilities for Records, you can nest and flatten Arrays using `flatten_arrays=true`:

```
{
  "__name__" : "Order",                                <-- artificial __name__ key for record
  "customer_id" : "111222",
  "order_details" : {                                  <-- array of records
    "0" : {                                             <-- array index 0
      "__name__" : "OrderDetail",
      "product_detail" : {
        "__name__" : "Product",
        "price" : "46.21",
        "product_category" : {                         <- array of strings
          "0" : "electronics",
          "1" : "printers",
          "2" : "computers"
        },
        "product_name" : "hp printer X11ew",
        "product_status" : "ONLY_FEW_LEFT"
      }
    },
    "order_id" : "2389646",
    "total" : "132.43"
  }
}
```

Here is the result of flattening the array:

```
{
  "0.order_details.__name__" : "OrderDetail",
  "0.order_details.product_detail.0.product_category" : "electronics",
  "0.order_details.product_detail.1.product_category" : "prnters",
  "0.order_details.product_detail.2.product_category" : "computers",
  "0.order_details.product_detail.__name__" : "Product",
  "0.order_details.product_detail.price" : "46.21",
  "0.order_details.product_detail.product_name" : "hp printer X11ew",
  "0.order_details.product_detail.product_status" : "ONLY_FEW_LEFT",
  "__name__" : "Order",
  "customer_id" : "111222",
  "order_id" : "2389646",
  "total" : "132.43"
}
```

## Maps

Maps use the type name map and support one attribute:

Attribute	Description
values	The schema of the map's items

The favroparser treats map keys as strings. For example, you can declare the map type as a long as follows:

```
{"type": "map", "values": "long"}
```

Similar to Records types, Maps can also be nested and flattened using `flatten_maps=true`.

The `favroparser` inserts key-value pairs from the Avro map as key-value pairs in the `__raw__` column. For an Avro record that has `KeyX` with value 10, and `KeyY` with value 20, `favroparser` loads the key-value pairs as virtual columns `KeyX` and `KeyY`, with values 10 and 20, respectively.

## Unions

Vertica uses JSON arrays to represent Avro Unions. Consider this example:

```
{"name": "TransactionID", "type": ["string", "null"]}
```

The field `TransactionID` can be a string or null.

## Fixed

Fixed (fixed) Avro types support two attributes:

Attribute	Description
name	A string for the name of this data type
size	An integer, specifying the number of bytes per value

For example, you can declare a 16-byte quantity:

```
{"type": "fixed", "size": 16, "name": "md5"}
```

With the preceding declaration is the Avro file schema, consider a record that contains a field with the following byte values for the key `md5`:

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 3, 4, 5]
```

The `favroparser` loads the key value pair as an `md5` key with the preceding byte values.

# Loading Common Event Format (CEF) Data

Use the flex parser `fcefparser` to load HPE ArcSight or other Common Event Format (CEF) log file data into columnar and flexible tables. For more information, see the [ArcSight Common Event Format \(CEF\) Guide](#).

When you use the parser to load arbitrary CEF-format files, it interprets key names in the data as virtual columns in your flex table. After loading, you can query your CEF data directly, regardless of which set of keys exist in each row. You can also use the associated flex table data and map functions to manage CEF data access.

## Create a Flex Table and Load CEF Data

This section uses a sample set of CEF data. All IP addresses have been purposely changed to be inaccurate, and Return characters added for illustration.

To use this sample data, copy the following text and remove all Return characters. Save the file as `CEF_sample.cef`, which is the name used throughout these examples.

```
CEF:0|ArcSight|ArcSight|6.0.3.6664.0|agent:030|Agent [test] type [testalertng] started|Low|
eventId=1 mrt=1396328238973 categorySignificance=/Normal categoryBehavior=/Execute/Start
categoryDeviceGroup=/Application catdt=Security Mangement categoryOutcome=/Success
categoryObject=/Host/Application/Service art=1396328241038 cat=/Agent/Started
deviceSeverity=Warning rt=1396328238937 fileType=Agent
cs2=<Resource ID\="3DxKlG0UBABCAA0cXXAZIwA\="/> c6a4=fe80:0:0:0:495d:cc3c:db1a:de71
cs2Label=Configuration Resource c6a4Label=Agent
IPv6 Address ahost=SKEELES10 agt=888.99.100.1 agentZoneURI=/All Zones/ArcSight
System/Private Address Space
Zones/RFC1918: 888.99.0.0-888.200.255.255 av=6.0.3.6664.0 atz=Australia/Sydney
aid=3DxKlG0UBABCAA0cXXAZIwA\=" at=testalertng dvchost=SKEELES10 dvc=888.99.100.1
deviceZoneURI=/All Zones/ArcSight System/Private Address Space Zones/RFC1918:
888.99.0.0-888.200.255.255 dtz=Australia/Sydney _cefVer=0.1
```

### 1. Create a flex table logs:

```
=> create flex table logs();
CREATE TABLE
```

### 2. Load the sample CEF file, using the flex parser fcefparser:

```
=> copy logs from '/home/dbadmin/data/CEF_sample.cef' parser fcefparser();
Rows Loaded
-----
1
(1 row)
```

### 3. Use the maptostring() function to see the contents of the logs flex table:

```
=> select maptostring(__raw__) from logs;
maptostring

-----
{
  "_cefver" : "0.1",
  "agentzoneuri" : "/All Zones/ArcSight System/Private Address
    Space Zones/RFC1918: 888.99.0.0-888.200.255.255",
  "agt" : "888.99.100.1",
  "ahost" : "SKEELES10",
  "aid" : "3DxKlG0UBABCAA0cXXAZIwA=",
  "art" : "1396328241038",
  "at" : "testalertng",
  "atz" : "Australia/Sydney",
```

```
"av" : "6.0.3.6664.0",
"c6a4" : "fe80:0:0:0:495d:cc3c:db1a:de71",
"c6a4label" : "Agent IPv6 Address",
"cat" : "/Agent/Started",
"catdt" : "Security Mangement",
"categorybehavior" : "/Execute/Start",
"categorydevicegroup" : "/Application",
"categoryobject" : "/Host/Application/Service",
"categoryoutcome" : "/Success",
"categorysignificance" : "/Normal",
"cs2" : "<Resource ID=\"3DxKlG0UBABCAA0cXXAZIwA==\"/>",
"cs2label" : "Configuration Resource",
"deviceproduct" : "ArcSight",
"deviceseverity" : "Warning",
"devicevendor" : "ArcSight",
"deviceversion" : "6.0.3.6664.0",
"devicezoneuri" : "/All Zones/ArcSight System/Private Address Space
    Zones/RFC1918: 888.99.0.0-888.200.255.255",
"dtz" : "Australia/Sydney",
"dvc" : "888.99.100.1",
"dvchost" : "SKEELES10",
"eventid" : "1",
"filetype" : "Agent",
"mrt" : "1396328238973",
"name" : "Agent [test] type [testalertng] started",
"rt" : "1396328238937",
"severity" : "Low",
"signatureid" : "agent:030",
"version" : "0"
}
(1 row)
```

## Create a Columnar Table and Load CEF Data

This example lets you compare the flex table for CEF data with a columnar table. You do so by creating a new table and load the same `CEF_sample.cef` file used in the preceding flex table example.

1. Create a columnar table, `col_logs`, defining the prefix names that are hard coded in `fcefparser`:

```
=> create table col_logs(version int,
    devicevendor varchar,
    deviceproduct varchar,
    deviceversion varchar,
    signatureid varchar,
    name varchar,
    severity varchar);
CREATE TABLE
```

2. Load the sample file into `col_logs`, as you did for the flex table:

```
=> copy col_logs from '/home/dbadmin/data/CEF_sample.cef' parser fcefparser();
Rows Loaded
-----
              1
(1 row)
```

3. Query the table. You can find the identical information in the flex table output.

```
=> \x
Expanded display is on.
VMart=> select * from col_logs;
-[ RECORD 1 ]-----
version      | 0
devicevendor | ArcSight
deviceproduct | ArcSight
deviceversion | 6.0.3.6664.0
signatureid  | agent:030
name         | Agent [test] type [testalertng] started
severity     | Low
```

## Compute Keys and Build a Flex Table View

In this example, you use a flex helper function to compute keys and build a view for the logs flex table.

1. Use the `compute_flextable_keys_and_build_view` function to compute keys and populate a view generated from the logs flex table:

```
=> select compute_flextable_keys_and_build_view('logs');
               compute_flextable_keys_and_build_view
-----
Please see public.logs_keys for updated keys
The view public.logs_view is ready for querying
(1 row)
```

2. Query the `logs_keys` table to see what the function computed from the sample CEF data:

```
=> select * from logs_keys;
 key_name      | frequency | data_type_guess
-----+-----+-----
c6a4           |          1 | varchar(60)
c6a4label     |          1 | varchar(36)
categoryobject |          1 | varchar(50)
categoryoutcome |          1 | varchar(20)
categorysignificance |          1 | varchar(20)
cs2           |          1 | varchar(84)
cs2label      |          1 | varchar(44)
```

deviceproduct	1	varchar(20)
deviceversion	1	varchar(24)
devicezoneuri	1	varchar(180)
dvchost	1	varchar(20)
version	1	varchar(20)
ahost	1	varchar(20)
art	1	varchar(26)
at	1	varchar(22)
cat	1	varchar(28)
catdt	1	varchar(36)
devicevendor	1	varchar(20)
dtz	1	varchar(32)
dvc	1	varchar(24)
filetype	1	varchar(20)
mrt	1	varchar(26)
_cefver	1	varchar(20)
agentzoneuri	1	varchar(180)
agt	1	varchar(24)
aid	1	varchar(50)
atz	1	varchar(32)
av	1	varchar(24)
categorybehavior	1	varchar(28)
categorydevicegroup	1	varchar(24)
deviceseverity	1	varchar(20)
eventid	1	varchar(20)
name	1	varchar(78)
rt	1	varchar(26)
severity	1	varchar(20)
signatureid	1	varchar(20)
(36 rows)		

### 3. Query several columns from the logs\_view:

```
=> \x
Expanded display is on.
VMart=> select version, devicevendor, deviceversion, name, severity, signatureid
        from logs_view;
-[ RECORD 1 ]-+-----
version      | 0
devicevendor | ArcSight
deviceversion | 6.0.3.6664.0
name         | Agent [test] type [testalertng] started
severity     | Low
signatureid  | agent:030
```

## Use the fcefpaser Delimiter Parameter

In this example, you use the `fcefpaser delimiter` parameter to query events located in California, New Mexico, and Arizona.



1. Create a new columnar table, CEFData3:

```
=> create table CEFData3(eventId int, location varchar(20));  
CREATE TABLE
```

2. Using the delimiter=' , ' parameter, load some CEF data into the table:

```
=> copy CEFData3 from stdin parser fcefparser(delimiter=',');  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> eventId=1,location=California  
>> eventId=2,location=New Mexico  
>> eventId=3,location=Arizona  
>> \.
```

3. Query the table:

```
=> select eventId, location from CEFData3;  
eventId | location  
-----+-----  
1 | California  
2 | New Mexico  
3 | Arizona  
(3 rows)
```

## Loading CSV Data

Use the `fcsvparser` to load data in CSV format (comma-separated values). Since no formal CSV standard exists, Vertica supports the [RFC 4180](#) standard as the default behavior for `fcsvparser`. Other parser parameters simplify various combinations of CSV options into columnar or flex tables. `fcsvparser` parses the following CSV data formats:

- **RFC 4180:** The RFC4180 csv format parser for Vertica flex tables. The parameters for this format are fixed and cannot be changed.
- **Traditional:** The traditional csv parser, which allows the user to specify the parameter values such as delimiter or record terminator. For a detailed list of parameters, please refer the [FCSVPARSER](#)

## Using Default Parser Settings

These fixed parameter settings apply to the RCF4180 format. You can change the default values for Traditional format type, if necessary.

Parameter	Data	Fixed	Default
-----------	------	-------	---------

	Type	Value (RCF4180)	Value (Traditional)
delimiter	CHAR	,	,
enclosed_by	CHAR	"	\
escape	CHAR	"	"
record_terminator	CHAR	\n; \r\n	\n; \r\n

Use the `type` parameter to indicate either an RFC 4180-compliant file or a traditional-compliant file. You can specify `type` as `RCF4180`. However, you must first verify that the data is compatible with the preceding fixed values for parameters of the RFC4180 format. The default value of the `type` parameter is `RCF4180`.

## Loading CSV Data (RFC4180)

Follow these steps to use `fcsvparser` to load data in the RFC4180 CSV data format.

To perform this task, assume that the CSV file to be loaded has the following sample contents:

```
$ more /home/dbadmin/flex/flexData1.csv
sno,name,age,gender
1,John,14,male
2,Mary,23,female
3,Mark,35,male
```

### 1. Create a flex table:

```
=> CREATE FLEX TABLE csv_basic();
CREATE TABLE
```

### 2. Load the data from csv file using `fcsvparser`:

```
=> COPY csv_basic FROM '/home/dbadmin/flex/flexData1.csv' PARSER fcsvparser();
Rows Loaded
-----
3(1 row)
```

### 3. View the data loaded in the flex table:

```
=> SELECT maptostring(__raw__) FROM csv_basic;
maptostring
-----
{
  "age" : "14",
  "gender" : "male",
```

```
"name" : "John",  
"sno" : "1"  
}  
{  
"age" : "23",  
"gender" : "female",  
"name" : "Mary",  
"sno" : "2"  
}  
{  
"age" : "35",  
"gender" : "male",  
"name" : "Mark",  
"sno" : "3"  
}  
(3 rows)
```

## Loading CSV Data (Traditional)

Follow these steps to use `fcsvparser` to load data in traditional CSV data format using `fcsvparser`.

In this example, the CSV file loaded has `delimiter` as `$` and `record_terminator` as `#`. Assume that the CSV file to be loaded has the following sample contents:

```
$ more /home/dbadmin/flex/flexData1.csv  
sno$name$age$gender#  
1$John$14$male#  
2$Mary$23$female#  
3$Mark$35$male#
```

### 1. Create a flex table:

```
=> CREATE FLEX TABLE csv_basic();  
CREATE TABLE
```

### 2. Load the data in flex table using `fscvparser` with parameters `type='traditional'`, `delimiter='$'` and `record_terminator='#'` :

```
=> COPY csv_basic FROM '/home/dbadmin/flex/flexData2.csv' PARSER fcsvparser  
(type='traditional',  
delimiter='$', record_terminator='#');  
Rows Loaded  
-----  
3  
(1 row)
```

### 3. View the data loaded in the flex table:

```
=> SELECT maptostring(__raw__) FROM csv_basic;  
maptostring  
-----
```

```
{
  "age" : "14",
  "gender" : "male",
  "name" : "John",
  "sno" : "1"
}
{
  "age" : "23",
  "gender" : "female",
  "name" : "Mary",
  "sno" : "2"
}
{
  "age" : "35",
  "gender" : "male",
  "name" : "Mark",
  "sno" : "3"
}
(3 rows)
```

**Note:** Refer to the following example to load new line characters using `fcsvparser`.

```
=> COPY foo_1 FROM STDIN PARSER fcsvparser(reject_on_materialized_type_error=true)
rejected data as table "m";
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> a,b
>> "10 <- we have quotes to escape the new line here
>> hi",20
>> 1,2
>> \.
```

## Rejecting on Duplicate Values

You can prevent loading duplicate data by using the `reject_on_duplicate=true` option with the `fcsvparser`. The next example shows how to use this parameter and then displays the specified exception and rejected data files.

**Note:** The `fcsvparser` rejects the entire load if any duplicate exists. The rejected data file includes the duplicate record that caused the load to fail. The exceptions file contains the reason for the rejection.

```
=> CREATE FLEX TABLE csv_basic();
CREATE TABLE

=> COPY csv_basic FROM stdin PARSER fdelimitedparser(reject_on_duplicate=true)
exceptions '/home/dbadmin/load_errors/except.out' rejected data '/home/dbadmin/load_
errors/reject.out';
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
```

```
>> A|A
>> 1|2
>> \.

=> \! cat /home/dbadmin/load_errors/reject.out
A|A
=> \! cat /home/dbadmin/load_errors/except.out
COPY: Input record 1 has been rejected (Processed a header row with duplicate keys with
reject_on_duplicate specified; rejecting.). Please see /home/dbadmin/load_errors/reject.out,
record 1 for the rejected record.
COPY: Loaded 0 rows, rejected 1 rows.
```

## Rejecting Data on Materialized Column Type Errors

The `fcsvparser` parser has a Boolean parameter, `reject_on_materialized_type_error`. Setting this parameter to `true` causes rows to be rejected if *both* the following conditions exist in the input data:

- Includes keys matching an existing materialized column
- Has a key value that cannot be coerced into the materialized column's data type

The following examples illustrate setting this parameter.

1. Create a table, `reject_true_false`, with two real columns:

```
=> CREATE FLEX TABLE reject_true_false(one int, two int);
CREATE TABLE
```

2. Load CSV data into the table (from STDIN), using the `fcsvparser` with `reject_on_materialized_type_error=false`. While `false` is the default value, you can specify it explicitly, as shown. Additionally, set the parameter `header=true` to specify the columns for input values:

```
=> COPY reject_true_false FROM stdin PARSER fcsvparser(reject_on_materialized_type_
error=false,header=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> one,two
>> 1,2
>> "3","four"
>> "five",6
>> 7,8
>> \.
```

3. Invoke `maptostring` to display the table values after loading data:

```
=> SELECT maptostring(__raw__), one, two FROM reject_true_false;
maptostring          | one | two
-----+-----+-----
{
"one" : "1",
"two" : "2"
}
|   1 |   2
{
"one" : "3",
"two" : "four"
}
|   3 |
{
"one" : "five",
"two" : "6"
}
|   |   6
{
"one" : "7",
"two" : "8"
}
|   7 |   8
(4 rows)
```

4. Truncate the table to empty the data stored in the table:

```
=> TRUNCATE TABLE reject_true_false;
TRUNCATE TABLE
```

5. Reload the same data again, but this time, set `reject_on_materialized_type_error=true`:

```
=> COPY reject_true_false FROM stdin PARSER fcsvparser(reject_on_materialized_type_
error=true,header=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> one,two
>> 1,2
>> "3","four"
>> "five",6
>> 7,8
>> \.
```

6. Call `maptostring` to display the table contents. Only two rows are currently loaded, whereas the previous results had four rows. The rows having input values with incorrect data type have been rejected:

```
=> SELECT maptostring(__raw__), one, two FROM reject_true_false;
maptostring          | one | two
-----+-----+-----
{
"one" : "1",
```

```
"two" : "2"  
}  
| 1 | 2  
{  
"one" : "7",  
"two" : "8"  
}  
| 7 | 8  
(2 rows)
```

**Note:** The parser `fcsvparser` uses null values if there is a type mismatch and you set the `reject_on_materialized_type_error` parameter to false.

## Rejecting or Omitting Empty Rows

Valid CSV files can include empty key and value pairs. Such rows are invalid for SQL. You can control the behavior for empty rows by either rejecting or omitting them, using two boolean `FCSVPARSER` parameters:

- `reject_on_empty_key`
- `omit_empty_keys`

The following example illustrates how to set these parameters:

1. Create a flex table:

```
=> CREATE FLEX TABLE csv_basic();  
CREATE TABLE
```

2. Load CSV data into the table (from STDIN), using the `fcsvparser` with `reject_on_empty_key=false`. While false is the default value, you can specify it explicitly, as shown. Additionally, set the parameter `header=true` to specify the columns for input values:

```
=> COPY csv_basic FROM stdin PARSER fcsvparser(reject_on_empty_key=false,header=true);  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> ,num  
>> 1,2  
>> \.
```

3. Invoke `maptostring` to display the table values after loading data:

```
=>SELECT maptostring(__raw__) FROM csv_basic;
```

```
maptostring
-----
{
  "" : "1",
  "num" : "2"
}
(1 row)
```

4. Truncate the table to empty the data stored in the table:

```
=> TRUNCATE TABLE csv_basic;
TRUNCATE TABLE
```

5. Reload the same data again, but this time, set `reject_on_empty_key=true`:

```
=> COPY csv_basic FROM stdin PARSE fcsvparser(reject_on_empty_key=true,header=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> ,num
>> 1,2
>> \.
```

6. Call `maptostring` to display the table contents. No rows are loaded because one of the keys is empty:

```
=>SELECT maptostring(__raw__) FROM csv_basic;
maptostring
-----
(0 rows)
```

7. Truncate the table to empty the data stored in the table:

```
=> TRUNCATE TABLE csv_basic;
TRUNCATE TABLE
```

8. Reload the same data again, but this time, set `omit_empty_keys=true`:

```
=> COPY csv_basic FROM stdin PARSE fcsvparser(omit_empty_keys=true,header=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> ,num
>> 1,2
>> \.
```

9. Call `maptostring` to display the table contents. One row is now loaded, and the rows with empty keys are omitted:

```
=> SELECT maptostring(__raw__) FROM csv_basic;
maptostring
```



```
-----  
{  
  "num" : "2"  
}  
(1 row)
```

**Note:** The `fcsvparser` uses 'col###' as a column name if no header row is indicated. If a table column name and key name matches, the parser loads the table column with values associated with the matching key.

## Using the NULL Parameter

Use the `NULL` parameter with `fcsvparser` to load `NULL` values into a flex table.

The next example uses this parameter:

1. Create a flex table.

```
=> CREATE FLEX TABLE fcsv(c1 int);  
CREATE TABLE
```

2. Load CSV data in flex table using `STDIN` and `NULL` parameter.

```
=> COPY fcsv FROM STDIN PARSER fcsvparser() delimiter '|' NULL 'NULL' ;  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> a,b,c1  
>> 10,20,NULL  
>> 20,30,50  
>> 20,30,40  
>> \.
```

3. Use `compute_flextable_keys_and_build_view` function to compute keys and build flex view.

```
=> SELECT compute_flextable_keys_and_build_view('fcsv');  
compute_flextable_keys_and_build_view  
-----  
  
-----  
Please see public.fcsv_keys for updated keys  
The view public.fcsv_view is ready for querying  
(1 row)
```

4. View the flex view and replace the `NULL` values.

```
=> SELECT * FROM public.fcsv_view;  
a | b | c1  
----+-----+----  
20 | 30 | 50
```

```
10 | 20 |  
20 | 30 | 40  
(3 rows)  
  
=> SELECT a,b, ISNULL(c1,-1) from public.fcsv_view;  
a | b | ISNULL  
-----  
20 | 30 | 50  
10 | 20 | -1  
20 | 30 | 40  
(3 rows)
```

## Loading Delimited Data

You can load flex tables with one of two delimited parsers, `fdelimitedparser` or `fdelimitedpairparser`.

- Use `fdelimitedpairparser` when the data specifies column names with the data in each row.
- Use `fdelimitedparser` when the data does not specify column names or has a header row for column names.

This section describes using some options that `fdelimitedpairparser` and `fdelimitedparser` support.

## Rejecting Duplicate Values

You can prevent loading duplicate data by using the `reject_on_duplicate=true` option with the `fdelimitedparser`. The next example shows how to use this parameter and then displays the specified exception and rejected data files.

**Note:** For the `fdelimitedparser`, the entire load is rejected if any duplicate is found. The exceptions file contains the reason for the rejection. The rejected data file includes the duplicate record that caused the load to fail.

```
=> create flex table delim_dupes();  
CREATE TABLE  
  
=> copy delim_dupes from stdin parser fdelimitedparser(reject_on_duplicate=true)  
exceptions '/home/dbadmin/load_errors/except.out' rejected data '/home/dbadmin/load_  
errors/reject.out';  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
  
>> A|A  
>> 1|2  
>> \.
```

```
=> \! cat /home/dbadmin/load_errors/reject.out
A|A
=> \! cat /home/dbadmin/load_errors/except.out
COPY: Input record 1 has been rejected (Processed a header row with duplicate keys with
reject_on_duplicate specified; rejecting.). Please see /home/dbadmin/load_errors/reject.out,
record 1 for the rejected record.
COPY: Loaded 0 rows, rejected 1 rows.
```

## Rejecting Materialized Column Type Errors

Both the `fjsonparser` and `fdelimitedparser` parsers have a boolean parameter, `reject_on_materialized_type_error`. Setting this parameter to `true` causes rows to be rejected if *both* the following conditions exist in the input data:

- Includes keys matching an existing materialized column
- Has a value that cannot be coerced into the materialized column's data type

Suppose the flex table has a materialized column, `OwnerPercent`, declared as a `FLOAT`. Trying to load a row with an `OwnerPercent` key that has a `VARCHAR` value causes `fdelimitedparser` to reject the data row.

The following examples illustrate setting this parameter.

1. Create a table, `reject_true_false`, with two real columns:

```
=> create flex table reject_true_false(one varchar, two int);
CREATE TABLE
```

2. Load JSON data into the table (from STDIN), using the `fjsonparser` with `reject_on_materialized_type_error=false`. While `false` is the default value, the following example specifies it explicitly for illustration:

```
=> copy reject_true_false from stdin parser fjsonparser(reject_on_materialized_type_
error=false);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> {"one": 1, "two": 2}
>> {"one": "one", "two": "two"}
>> {"one": "one", "two": 2}
>> \.
```

3. Invoke `maptostring` to display the table values after loading data:

```
=> select maptostring(__raw__), one, two from reject_true_false;
      maptostring      | one | two
-----+-----+-----
{
```

```
    "one" : "one",
    "two" : "2"
  }
  | one | 2
{
  "one" : "1",
  "two" : "2"
}
  | 1 | 2
{
  "one" : "one",
  "two" : "two"
}
  | one |
(3 rows)
```

4. Truncate the table:

```
=> truncate table reject_true_false;
```

5. Reload the same data again, but this time, set `reject_on_materialized_type_error=true`:

```
=> copy reject_true_false from stdin parser fjsonparser(reject_on_materialized_type_
error=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> {"one": 1, "two": 2}
>> {"one": "one", "two": "two"}
>> {"one": "one", "two": 2}
>> \.
```

6. Call `maptostring` to display the table contents. Only two rows were loaded, whereas the previous results had three rows:

```
=> select maptostring(__raw__), one, two from reject_true_false;
      maptostring      | one | two
-----+-----+-----
{
  "one" : "1",
  "two" : "2"
}
  | 1 | 2
{
  "one" : "one",
  "two" : "2"
}
  | one | 2
(2 rows)
```

## Loading JSON Data

You can load JSON data into flex or columnar tables. This section describes some examples of using the `fjsonparser` with several of the parser options.

## Checking JSON Integrity

Before loading any JSON data, be sure that the data is valid. You can verify JSON data integrity using a web tool such as [JSONLint](#). Copy your JSON data into the tool. If any data is invalid, the tool returns a message similar to the one in this example:

```
Parse error on line 170:...257914002502451200}{    "id_str": "257
-----^
Expecting 'EOF', '}', ',', ']', ']'
```

## Using `flatten_maps` and `flatten_arrays` Parameters

When loading JSON data, the `fjsonparser` uses the parameters `flatten_maps` and `flatten_arrays` to control how the parser handles the data it is loading. Here are the default settings for these two parameters:

Parameter	Default	Change Default
<code>flatten_maps</code>	TRUE: Flatten all maps.	<code>flatten_maps=FALSE</code>
<code>flatten_arrays</code>	FALSE: Do not flatten arrays.	<code>flatten_arrays=TRUE</code>

You control the default the behavior by using one or both `flatten_` parameters.

For JSON maps, the parser flattens all submaps, separating the levels with a period (.). Consider the following input data with a submap:

```
{ grade: { level: 4 } }
```

The default parser behavior results in the following map:

```
{ "grade.level" -> "4" }
```

For JSON arrays, the parser maintains the array. Consider the following input data containing a 2-element array, with values 1 and 2:

```
{ grade: [ 1 2 ] }
```

The default parser behavior results in the following array:

```
{ "grade": { "0" -> "1", "1" -> "2" } }
```

**Note:** Using the parameters `flatten_maps` and `flatten_arrays` is recursive, and affects all data.

## Loading from a Specific Start Point

You can use the `fjsonparser start_point` parameter to load JSON data beginning at a specific key, rather than at the beginning of a file. Data is parsed from after the `start_point` key until the end of the file, or to the end of the first `start_point`'s value. The `fjsonparser` ignores any subsequent instance of the `start_point`, even if that key appears multiple times in the input file. If the input data contains only one copy of the `start_point` key, and that value is a list of JSON elements, the parser loads each element in the list as a row.

This section uses the following sample JSON data, saved to a file (`alphanums.json`):

```
{ "A": { "B": { "C": [ { "d": 1, "e": 2, "f": 3 }, { "g": 4, "h": 5, "i": 6 },  
  { "j": 7, "k": 8, "l": 9 } ] } } }
```

1. Create a flex table, `start_json`:

```
=> create flex table start_json();  
CREATE TABLE
```

2. Load `alphanums.json` into `start_json` using the `fjsonparser` without any parameters:

```
=> copy start_json from '/home/dbadmin/data/flex/alphanums.json' parser fjsonparser();  
Rows Loaded  
-----  
1  
(1 row)
```

3. Use `maptostring` to see the results of loading all of `alphanums.json`:

```
=> select maptostring(__raw__) from start_json;  
maptostring  
-----  
{  
  "A.B.C" : {  
    "0.d" : "1",  
    "0.e" : "2",  
    "0.f" : "3",  
    "1.g" : "4",
```

```
      "1.h" : "5",  
      "1.i" : "6",  
      "2.j" : "7",  
      "2.k" : "8",  
      "2.l" : "9"  
    }  
  }  
(1 row)
```

4. Truncate start\_json and load alphanums.json with the start\_point parameter:

```
=> truncate table start_json;  
TRUNCATE TABLE  
=> copy start_json from '/home/dbadmin/data/flex/alphanums.json' parser  
-> fjsonparser(start_point='B');  
Rows Loaded  
-----  
1  
(1 row)
```

5. Next, call maptostring again to compare the results of loading alphanums.json from start\_point='B':

```
=> select maptostring(__raw__) from start_json;  
      maptostring  
-----  
{  
  "C" : {  
    "0.d" : "1",  
    "0.e" : "2",  
    "0.f" : "3",  
    "1.g" : "4",  
    "1.h" : "5",  
    "1.i" : "6",  
    "2.j" : "7",  
    "2.k" : "8",  
    "2.l" : "9"  
  }  
}  
(1 row)
```

## Parsing From a Start Point Occurrence

If a start\_point value occurs in multiple locations in your JSON data, you can use the start\_point\_occurrence integer parameter to specify the occurrence at which to start parsing. By defining start\_point\_occurrence, fjsonparser begins at the nth occurrence of start\_point.

## Controlling Column Name Separators

By default, `fjsonparser` produces column names by concatenating JSON field names with a period (.). You can change the default separator by specifying a different character with the `key_separator` parameter.

## Handling Special Characters

Some input JSON data can have special characters in field names. You can replace these characters by setting the `suppress_nonalphanumeric_key_chars` to `TRUE`. With this parameter setting, all special characters are converted to an underscore (\_) character.

## Rejecting on Duplicate Values

You can avoid loading duplicate data by using the `reject_on_duplicate=true` option with the `fjsonparser`. The next example uses this option while loading data and then displays the specified exception and rejected data files.

**Note:** For the `fjsonparser`, the entire load is rejected if any duplicate is found. To save disk space, rejected data is not saved. The rejected data file includes one newline per rejected record, and the exceptions file includes the reason for the rejection.

```
=> create flex table json_dupes();
CREATE TABLE

=> copy json_dupes from stdin parser fjsonparser(reject_on_duplicate=true)
exceptions '/home/dbadmin/load_errors/json_e.out'
rejected data '/home/dbadmin/load_errors/json_r.out';

Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.

>> {"a":"1","a":"2","b":"3"}
>> \.

=> \!cat /home/dbadmin/load_errors/json_e.out
COPY: Input record 1 has been rejected (Rejected by user-defined parser).
Please see /home/dbadmin/load_errors/json_r.out, record 1 for the rejected record.
COPY: Loaded 0 rows, rejected 1 rows.
```

## Rejecting Data on Materialized Column Type Errors

Both the `fjsonparser` and `fdelimitedparser` parsers have a Boolean parameter, `reject_on_materialized_type_error`. Setting this parameter to `true` causes rows to be rejected if the input data:



- Includes keys matching an existing materialized column
- Has a key value that cannot be coerced into the materialized column's data type.

The following examples illustrate setting this parameter.

1. Create a table, `reject_true_false`, with two real columns:

```
=> create flex table reject_true_false(one varchar, two int);  
CREATE TABLE
```

2. Load JSON data into the table (from STDIN), using the `fjsonparser` with `reject_on_materialized_type_error=false`. While `false` is the default value, the following example specifies it explicitly for illustration:

```
=> copy reject_true_false from stdin parser  
-> fjsonparser(reject_on_materialized_type_error=false);  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> {"one": 1, "two": 2}  
>> {"one": "one", "two": "two"}  
>> {"one": "one", "two": 2}  
>> \.
```

3. Invoke `maptostring` to display the table values after loading data:

```
=> select maptostring(__raw__), one, two from reject_true_false;  
          maptostring      | one | two  
-----+-----+-----  
{  
  "one" : "one",  
  "two" : "2"  
}  
| one |  2  
{  
  "one" : "1",  
  "two" : "2"  
}  
| 1   |  2  
{  
  "one" : "one",  
  "two" : "two"  
}  
| one |  
(3 rows)
```

4. Truncate the table:

```
=> truncate table reject_true_false;
```

5. Reload the same data again, but this time, set `reject_on_materialized_type_error=true`:

```
=> copy reject_true_false from stdin parser fjsonparser(reject_on_materialized_type_error=true);
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> {"one": 1, "two": 2}
>> {"one": "one", "two": "two"}
>> {"one": "one", "two": 2}
>> \.
```

6. Call `maptostring` to display the table contents. Only two rows were loaded, whereas the previous results had three rows:

```
=> select maptostring(__raw__), one, two from reject_true_false;
      maptostring      | one | two
-----+-----+-----
{
  "one" : "1",
  "two" : "2"
}
| 1 | 2
{
  "one" : "one",
  "two" : "2"
}
| one | 2
(2 rows)
```

## Rejecting or Omitting Empty Rows

Valid JSON files HPE can include empty key and value pairs, such as this one:

```
{"": 1 "}
```

Such rows are invalid for SQL. To prevent this situation, you can control the behavior for empty rows, either rejecting or omitting them. You do so using two boolean parameters for the parsers `FDELIMITEDPARSER` or `FJSONPARSER`:

- `reject_on_empty_key`
- `omit_empty_keys`

## See Also

- [Using COPY with Kafka](#)

## Loading Matches from Regular Expressions

You can load flex or columnar tables with the matched results of a regular expression, using the `fregexparser`. This section describes some examples of using the options that the flex parsers support.

### Sample Regular Expression

These examples use the following regular expression, which searches information that includes the `timestamp`, `date`, `thread_name`, and `thread_id`. For illustrative purposes, this regular expression uses new lines to break up long text. Remove any new line characters before using this example in your own tests.

```
'^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)  
(?<thread_name>[A-Za-z ]+):(?<thread_id>\d{0-9a-f}+)-?(?<transaction_id>[0-9a-f]+)?  
(?:\[ (?<component>\w+)\] \<(?<level>\w+)\> )?(?:\<(?<level>\w+)\> @\[ (?<enode>\w+)\]?: )?  
(?<text>.*).'
```

## Using Regular Expression Matches for a Flex Table

You can load the results from a regular expression into a flex table, using the `fregexparser`. For a complete example of doing so, see [FREGEXPARSER](#).

### Using `fregexparser` for Columnar Tables

This section illustrates how to load the results of a regular expression used against a sample log file for a Vertica database. By using an external table definition, the This section presents an example of using the `fregexparser` to load data into a columnar table. Using a flex table parser for a columnar tables gives you the capability to mix data loads in one table, such as loading the results of a regular expression in one session, and JSON data in another.

The following basic examples illustrate this usage.

1. Create a columnar table, `vlog`, with the following columns:

```
=> CREATE TABLE vlog (  
    "text"                varchar(2322),  
    thread_id             varchar(28),  
    thread_name          varchar(44),  
    "time"                varchar(46),  
    component             varchar(30),
```

```
level          varchar(20),
transaction_id varchar(32),
elevel         varchar(20),
enode          varchar(34)
);
```

2. Use COPY to load parts of a log file using the sample regular expression, with the fregexparser:

```
=> COPY v_log FROM '/home/dbadmin/data/flex/vertica.log' PARSE
FRegexParser(pattern='^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)
(?<thread_name>[A-Za-z ]+):(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?
(?:\[(?<component>\w+)\] <(?<level>\w+)> )?(?:<(?<elevel>\w+)> @[?<enode>\w+)\]? : )
?(?<text>.*)') rejected data as table fregex_reject;
```

3. Query the time column:

```
=> select time from flogs limit 10;
      time
-----
2014-04-02 04:02:02.613
2014-04-02 04:02:02.613
2014-04-02 04:02:02.614
2014-04-02 04:02:51.008
2014-04-02 04:02:51.010
2014-04-02 04:02:51.012
2014-04-02 04:02:51.012
2014-04-02 04:02:51.013
2014-04-02 04:02:51.014
2014-04-02 04:02:51.017
(10 rows)
```

## Using External Tables with fregexparser

By creating an external columnar table for your Vertica log file, querying the table will return updated log information. The following basic example illustrate this usage.

1. Create a columnar table, `vertica_log`, using the `AS COPY` clause and `fregexparser` to load matched results from the regular expression. For illustrative purposes, the regular expression has new lines to break up long text:

```
=> CREATE EXTERNAL TABLE public.vertica_log
(
  "text" varchar(2322),
  thread_id varchar(28),
  thread_name varchar(44),
  "time" varchar(46),
  component varchar(30),
  level varchar(20),
  transaction_id varchar(32),
```

```

        elevel varchar(20),
        enode varchar(34)
    )
AS COPY
FROM '/home/dbadmin/data/vertica.log'
PARSER FRegexParser(pattern='^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)
(?<thread_name>[A-Za-z ]+):(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?
(?:\[?(?<component>\w+)\] \<(?(?<level>\w+)\> )?(?:<(?(?<elevel>\w+)\> @\[?(?<enode>\w+)\]\?: )?
(?(?<text>.*))');

```

## 2. Query from the external table to get updated results:

```

=> select component, thread_id, time from vertica_log limit 10;

```

component	thread_id	time
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
Init	0x16321430	2014-04-02 04:02:02.613
	0x16321430	2014-04-02 04:02:02.614
	0x16321430	2014-04-02 04:02:02.614
	0x16321430	2014-04-02 04:02:02.614

(10 rows)

## Computing Flex Table Keys

After loading data into a flex table, you must determine what key value pairs exist as populated virtual columns in the data. Two helper functions compute keys from map data:

- [COMPUTE\\_FLEXTABLE\\_KEYS](#)— See also [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)— Performs the same functionality as [COMPUTE\\_FLEXTABLE\\_KEYS](#) but also builds a new view. See also [Updating Flex Table Views](#).

## Using COMPUTE\_FLEXTABLE\_KEYS

Call this function with a flex table argument to compute a list of keys from the map data:

Calling the Function	Results
<code>compute_flextable_keys ('flex_table')</code>	Computes keys from the <i>flex_table</i> map data and populates the associated <i>flex_table_keys</i> with the virtual columns.

## Calculating Key Value Column Widths

During execution, this function determines a data type for each virtual column. It casts the values computed to VARCHAR, LONG VARCHAR, or LONG VARBINARY. Casting choice depends on the length of the key and whether the key includes nested maps.

The following examples illustrate this function. It shows results of populating the `_keys` table, after you create a flex table (`darkdata1`) and load data:

```
=> create flex table darkdata1();
CREATE TABLE
=> copy darkdata1 from '/test/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
      12
(1 row)
=> select compute_flextable_keys('darkdata1');
      compute_flextable_keys
-----
Please see public.darkdata1_keys for updated keys
(1 row)
=> select * from darkdata1_keys;
      key_name                                     | frequency | data_type_guess
-----+-----+-----
```

```

contributors          |      8 | varchar(20)
coordinates           |      8 | varchar(20)
created_at            |      8 | varchar(60)
entities.hashtags     |      8 | long varbinary(186)
entities.urls         |      8 | long varbinary(32)
entities.user_mentions |      8 | long varbinary(674)
.
.
.
retweeted_status.user.time_zone |      1 | varchar(20)
retweeted_status.user.url      |      1 | varchar(68)
retweeted_status.user.utc_offset |      1 | varchar(20)
retweeted_status.user.verified |      1 | varchar(20)
(125 rows)

```

The flex keys table has these columns:

Column	Description
key_name	The name of the virtual column (key).
frequency	The number of times the virtual column occurs in the map.
data_type_guess	<p>The data type for each virtual column. This value is cast to VARCHAR, LONG VARCHAR or LONG VARBINARY. Casting depends on the length of the key and whether the key includes one or more nested maps.</p> <p>In the <code>_keys</code> table output, the <code>data_type_guess</code> column values are also followed by a value in parentheses, such as <code>varchar(20)</code>. The value indicates the padded width of the key column. The width is determined by calculating the longest field, multiplied by the <code>FlexTableDataTypeGuessMultiplier</code> configuration parameter value. For more information, see <a href="#">Setting Flex Table Parameters</a>.</p>

# Materializing Flex Tables

Once flex tables exist, you can change the table structure to promote virtual columns to materialized (real) columns. If your table is already a hybrid table, you can change existing real columns and promote other important virtual columns. This section describes some key aspects of promoting columns, adding columns, specifying constraints, and declaring default values. It also presents some differences when loading flex or hybrid tables, compared with columnar tables.

**Note:** Materializing virtual columns by promoting them to real columns can significantly improve query performance. Vertica recommends that you materialize important virtual columns before running large and complex queries. Promoted columns cause a small increase in load performance.

## Adding Columns to Flex Tables

Add columns to your flex tables to promote virtual columns:

1. Add a column with the same name as a map key:

```
=> alter table darkdata1 add column "user.name" varchar;  
ALTER TABLE
```

2. Loading data into a materialized column populates the new column automatically:

```
=> copy darkdata1 from '/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();  
Rows Loaded  
-----  
12  
(1 row)
```

3. Query the materialized column from the flex table:

```
=> select "user.name" from darkdata1;  
user.name  
-----  
I'm Toasterâ  
Flu Beach  
seydo shi  
The End  
Uptown gentleman.  
~G A B R I E L A â  
Avita Desai  
laughing at clouds.  
(12 rows)
```



## Adding Columns with Default Values

The section [Using COPY with Flex Tables](#) describes the use of default values, and how they are evaluated during loading. As with all tables, using COPY to load data ignores any column default values.

**Note:** Adding a table column default expression to a flex table requires casting the column to an explicit data type.

1. Create a `darkdata1` table with some column definition, but a name that does not correspond to any key names in the JSON data you'll load. Assign a default value for a column you know exists in your data (`"user.lang"`):

```
=> create flex table darkdata1(talker long varchar default "user.lang");  
CREATE TABLE
```

2. Load some JSON data:

```
=> copy darkdata1 from '/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();  
Rows Loaded  
-----  
12  
(1 row)
```

3. Query the `talker` column to see that the default value was not used. The column contains NULL values.
4. Load data again, specifying just the `__raw__` column to use the column's default value:

```
=> copy darkdata1 (__raw__)from '/test/vertica/flextable/DATA/tweets_12.json'  
parser fjsonparser();  
Rows Loaded  
-----  
12  
(1 row)
```

5. Query to see that the column's default expression was used (`"user.lang"`), because you specified `__raw__`:

```
=> select "talker" from darkdata1;  
talker  
-----  
it
```

```
en
es
en
en
es
tr
en
(12 rows)
```

6. Alter the table to add a row with a key value name, assigning the key name as the default value (recommended):

```
=> alter table darkdata1 add column "user.name" varchar default "user.name";
ALTER TABLE
```

7. Load data again, this time without `__raw__`:

```
=> copy darkdata1 from '/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();
```

8. Query the two real columns and see that `talker` is NULL, since you did not specify the `__raw__` column. The `user.lang` column contains values from the data you loaded:

```
=> select "talker", "user.name" from darkdata1;
talker |      user.name
-----+-----
      | laughing at clouds.
      | Avita Desai
      | I'm Toasterâ
      |
      |
      | Uptown gentleman.
      | ~G A B R I E L A â
      | Flu Beach
      |
      | seydo shi
      | The End
(12 rows)
```

9. Load data once more, this time specifying a COPY statement default value expression for `user.name`:

```
=> copy darkdata1 (__raw__, "user.name" as 'QueenElizabeth'::varchar) from
'/test/vertica/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
      12
```

```
(1 row)
```

10. Query once more. `talker` has its default values (you used `__raw__`), and the `COPY` value expression (`QueenElizabeth`) overrode the `user.name` default column value:

```
=> select "talker", "user.name" from darkdata1;
talker |  user.name
-----+-----
it      | QueenElizabeth
en      | QueenElizabeth
es      | QueenElizabeth
        | QueenElizabeth
        | QueenElizabeth
        | QueenElizabeth
en      | QueenElizabeth
en      | QueenElizabeth
es      | QueenElizabeth
        | QueenElizabeth
tr      | QueenElizabeth
en      | QueenElizabeth
(12 rows)
```

To summarize, you can set a default column value as part of the `ALTER TABLE...ADD COLUMN...` operation. For materializing columns, the default should reference the key name of the virtual column (as in `"user.lang"`). Subsequently loading data with a `COPY` value expression overrides the default value of the column definition.

## Changing the `__raw__` Column Size

You can change the default size of the `__raw__` column for flex tables you plan to create, the current size of an existing flex table, or both.

To change the default size for the flex table `__raw__` column, use the following configuration parameter (described in [Setting Flex Table Parameters](#)):

```
=> ALTER DATABASE mydb SET FlexTableRowSize = 120000;
```

Changing the configuration parameter affects all flex tables you create after making this change.

To change the size of the `__raw__` column in an existing flex table, use the [ALTER TABLE](#) statement as follows:

```
=> alter table tester alter column __raw__ set data type long varbinary(120000);
ALTER TABLE
```

**Note:** An error will occur if you try reducing the `__raw__` column size to a value smaller than the data that the column already contains.

## Changing Flex Table Real Columns

You can make the following changes to the flex table real columns (`__raw__` and `__identity__`), but not to any virtual columns:

Actions	<code>__raw__</code>	<code>__identity__</code>
Change NOT NULL constraints (default)	Yes	Yes
Add primary key and foreign key (PK/FK) constraints	No	Yes
Create projections	No	Yes
Segment	No	Yes
Partition	No	Yes
Specify a user-defined scalar function (UDSF) as a default column expression in <code>ALTER TABLE x ADD COLUMN y</code> statement	No	No

**Note:** While segmenting and partitioning the `__raw__` column is permitted, it is not recommended due to its long data type. By default, if no real columns exist, flex tables are segmented on the `__identity__` column.

## Dropping Flex Table Columns

There are two considerations about dropping columns:

- You cannot drop the last column in your flex table's sort order.
- If you have not created a flex table with any real columns, or materialized any columns, you cannot drop the `__identity__` column.

## Updating Flex Table Views

Creating a flex table also creates a default view to accompany the table. The view has the name of the table with an underscore (`_view`) suffix. When you perform a `select` query from the default view, you get a prompt to run the helper function, as shown in this example:

```
=> \dv dark*

      List of View Fields
Schema |      View      | Column |      Type      | Size
-----+-----+-----+-----+-----
public | darkdata_view  | status | varchar(124)   | 124
public | darkdata1_view | status | varchar(124)   | 124
(2 rows)

=> select * from darkdata_view;

                                     status
-----
Please run compute_flextable_keys_and_build_view() to update this view to reflect real and
virtual columns in the flex table
(1 row)
```

Two helper functions create views:

- [COMPUTE\\_FLEXTABLE\\_KEYS](#)— See also [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)— Performs the same functionality as [BUILD\\_FLEXTABLE\\_KEYS](#) but also computes keys. See also [Using COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#).

## Using BUILD\_FLEXTABLE\_VIEW

After the function computes keys for the Flex Table ([Computing Flex Table Keys](#)), call this function with one or more arguments. The records under the `key_name` column of the `{flextable}_keys` table are used as view columns, along with any values for the key. If no values exist, the column value is NULL.

Regardless of the number of arguments, calling this function replaces the contents of the existing view as follows:

Function Invocation	Results
<code>build_flextable_view ('flexible_table')</code>	Changes the existing view associated with <i>flexible_table</i> with the current contents of the associated <i>flexible_table_keys</i> table.

Function Invocation	Results
<code>build_flextable_view ('flexible_table', 'view_name')</code>	Changes the view you specify with <i>view_name</i> by using the current contents of the <i>{flextable}_keys</i> table.
<code>build_flextable_view ('flexible_table', 'view_name', 'table_keys')</code>	Changes the view you specify with <i>view_name</i> to the current contents of the <i>flexible_table_keys</i> table. Use this function to change a view of your choice with the contents of the keys of interest.

If you do not specify a *view\_name* argument, the default name is the flex table name with a *\_view* suffix. For example, if you specify the table *darkdata* as the sole argument to this function, the default view is called *darkdata\_view*.

You cannot specify a custom view name with the same name as the default view *flex\_table\_view*, unless you first drop the default-named view and then create your own view of the same name.

Creating a view stores a definition of the column structure at the time of creation. Thus, if you create a flex table view and then promote virtual columns to real columns, you must rebuild the view. Querying a rebuilt flex table view that has newly promoted real columns produces two results. These results reflect values from both virtual columns in the map data and real columns.

## Handling JSON Duplicate Key Names in Views

SQL is a case-insensitive language, so the names *TEST*, *test*, and *TeSt* are identical. JSON data is case sensitive, so that it can validly contain key names of different cases with separate values.

When you build a flex table view, the function generates a warning if it detects same-name keys with different cases in the *{flextable}\_keys* table. For example, calling `BUILD_FLEXTABLE_VIEW` or `COMPUTE_FLEXTABLE_KEYS_AND_BUILD_VIEW()` on a flex table with duplicate key names results in these warnings:

```
=> select compute_flextable_keys_and_build_view('dupe');
```

```
WARNING 5821: Detected keys sharing the same case-insensitive key name
WARNING 5909: Found and ignored keys with names longer than the maximum column-name length limit
```

```
compute_flextable_keys_and_build_view
```

---

Please see `public.dupe_keys` for updated keys  
The view `public.dupe_view` is ready for querying  
(1 row)

While a `{flextable}_keys` table can include duplicate key names with different cases, a view cannot. Creating a flex table view with either of the helper functions consolidates any duplicate key names to one column name, consisting of all lowercase characters. All duplicate key values for that column are saved. For example, if these key names exist in a flex table:

- test
- Test
- tEst

The view will include a virtual column `test` with values from the `test`, `Test`, and `tEst` keys.

**Note:** The examples in this section include added Return characters to reduce line lengths. The product output may differ.

For example, consider the following query, showing the duplicate test key names:

[illegible]

```
-----+-----
key_name      | TEST
frequency     | 8
data_type_guess | varchar(20)
-[ RECORD 5 ]--+-----
key_name      | TEST
frequency     | 8
data_type_guess | varchar(20)
```

The following query displays the dupe flex table (dupe\_view). It shows the consolidated test and testtesttest... virtual columns. All the test, Test, and tEst virtual column values are in the test column:

```
=> select * from dupe_view;  
   test |  
testtesttesttesttesttesttesttesttesttesttesttesttesttesttesttesttesttesttest  
testtesttesttesttesttesttesttest  
-----+-----  
upper2 |  
half4  |  
lower1 |  
upper1 |  
half1  |  
half4  |  
       |  
lower1 |  
half1  |  
upper2 |  
       |  
lower2 |  
lower3 |  
upper1 |  
lower2 |  
lower3 |  
(16 rows)
```

## Creating a Flex Table View

The following example shows how to create a view, `dd_view`, from the flex table `darkdata`, which contains JSON data.

```
=> create view dd_view as select "user.lang"::varchar, "user.name"::varchar from darkdata;
CREATE VIEW
```

Query the key names you specified, and their values:

```
=> select * from dd_view;
  user.lang |      user.name
-----+-----
  en        | Uptown gentleman.
  en        | The End
```



## Using Flex Tables

### Updating Flex Table Views

```
it      | laughing at clouds.
es      | I'm Toasterâ
        |
en      | ~G A B R I E L A â
        |
en      | Avita Desai
tr      | seydo shi
        |
es      | Flu Beach
(12 rows)
```

This example shows how to call `build_flextable_view` with the original table and the view you previously created, `dd_view`:

```
=> select build_flextable_view ('darkdata', 'dd_view');
        build_flextable_view
-----
The view public.dd_view is ready for querying
(1 row)
```

Query the view again. You can see that the function populated the view with the contents of the `darkdata_keys` table. Next, review a snippet from the results, with the `key_name` columns and their values:

```
=> \x
Expanded display is on.
```

```
=> select * from dd_view;
.
.
.
user.following          |
user.friends_count      | 791
user.geo_enabled         | F
user.id                  | 164464905
user.id_str              | 164464905
user.is_translator       | F
user.lang                | en
user.listed_count        | 4
user.location             | Uptown..
user.name                 | Uptown gentleman.
.
.
.
```

When building views, be aware that creating a view stores a definition of the column structure at the time the view is created. If you promote virtual columns to real columns after building a view, the existing view definition is not changed. Querying this view with a select statement such as the following, returns values from only the `__raw__` column:

```
=> select * from myflextable_view;
```

Also understand that rebuilding the view after promoting virtual columns changes the resulting value. Future queries return values from both virtual columns in the map data and from real columns.

## Using COMPUTE\_FLEXTABLE\_KEYS\_AND\_BUILD\_VIEW

Call this function with a flex table to compute Flex table keys (see [Computing Flex Table Keys](#) ), and create a view in one step.

## Querying Flex Tables

After you create your flex table (with or without additional columns) and load data, you can perform four types of queries:

- SELECT
- COPY
- TRUNCATE
- DELETE

You can use SELECT queries for virtual columns that exist in the `__raw__` column and real columns in your flex tables. Column names are case insensitive.

## Unsupported DDL and DML Statements

You cannot use the following DDL and DML statements with flex tables:

```
CREATE TABLE flex_table AS...CREATE TABLE flex_table LIKE...  
SELECT INTO  
UPDATE  
MERGE  
INSERT INTO
```

## Querying Flex Table Keys

If you reference an undefined column ('`which_column`') in a flex table query, Vertica converts the query to a call to the `maplookup()` function as follows:

```
maplookup(__raw__, 'which_column')
```

The `maplookup()` function searches the VMap data for the requested key and returns the following information:

- String values associated with the key for a row.
- NULL if the key is not found.

For more information about handling NULL values, see [MAPCONTAINSKEY\(\)](#).

## Determining Flex Table Data Contents

If you don't know what your flex table contains, two helper functions let you explore that data to determine contents. Use these functions to compute the keys in the flex table `__raw__` column and, optionally, build a view based on those keys:

- [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)

For more information about these and other helper functions, see [Flex Data Functions Reference](#)

To determine what virtual columns exist:

1. Call the function as follows:

```
=> select compute_flextable_keys('darkdata');
compute_flextable_keys
-----
Please see public.darkdata_keys for updated keys(1 row)
```

2. View the computed key names by querying the `darkdata_keys` table:

```
=> select * from darkdata_keys;
```

key_name	frequency	data_type_guess
contributors	8	varchar(20)
coordinates	8	varchar(20)
created_at	8	varchar(60)
entities.hashtags	8	long varbinary(186)
.		
.		
retweeted_status.user.time_zone	1	varchar(20)
retweeted_status.user.url	1	varchar(68)
retweeted_status.user.utc_offset	1	varchar(20)
retweeted_status.user.verified	1	varchar(20)

(125 rows)

## Querying Virtual Columns

Continuing with the JSON data example, use select queries to explore content from the virtual columns. Then, analyze what's most important to you. This example shows querying some common virtual columns in the map data:

```
=> select "user.name", "user.lang", "user.geo_enabled" from darkdata1;
      user.name | user.lang | user.geo_enabled
-----+-----+-----
laughing at clouds. | it      | T
Avita Desai       | en      | F
I'm Toasterâ    | es      | T
                  |         |
                  |         |
Uptown gentleman. | en      | F
~G A B R I E L A â | en      | F
Flu Beach         | es      | F
                  |         |
seydo shi         | tr      | T
The End           | en      | F
(12 rows)
```

## Using Functions and Casting in Flex Table Queries

You can cast the virtual columns as required and use functions in your select queries. The next example queries the darkdata1 flex table for the created\_at and retweet\_count virtual columns, casting their values in the process:

```
=> select "created_at"::TIMESTAMP, "retweet_count"::INT from darkdata1 order by 1 desc;
      created_at | retweet_count
-----+-----
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:05 | 0
2012-10-15 14:41:04 | 1
                  |
                  |
                  |
(12 rows)
```

The following query uses the COUNT and AVG functions to determine the average length of text in different languages:

```
=> select "user.lang", count (*), avg(length("text"))::int from darkdata1 group by 1 order by 2 desc;
      user.lang | count | avg
-----+-----+-----
en             | 4     | 42
               | 4     |
es             | 2     | 96
it             | 1     | 50
```

```
tr      |      1 | 16  
(5 rows)
```

## Casting Data Types in a Query

The following query requests the values of the `created_at` virtual column, without casting to a specific data type:

```
=> select "created_at" from darkdata1;  
      created_at  
-----  
Mon Oct 15 18:41:04 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
Mon Oct 15 18:41:05 +0000 2012  
(12 rows)
```

The next example queries the same virtual column, casting `created_at` to a `TIMESTAMP`. Casting results in different output and the regional time:

```
=> select "created_at"::TIMESTAMP from darkdata1 order by 1 desc;  
      created_at  
-----  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:05  
2012-10-15 14:41:04
```

## Accessing an Epoch Key

The term *EPOCH* (all uppercase letters) is reserved in Vertica for internal use.

If your JSON data includes a virtual column called `epoch`, you can query it within your flex table. However, use the `maplookup()` function to do so.

## Querying Flex Views

Flex tables offer the ability of dynamic schema through the application of query rewriting. Use flex views to support restricted access to flex tables. As with flex tables, each time you use a `select` query on a flex table view, internally, Vertica invokes the `maplookup()` function, to return information on all virtual columns. This query behavior occurs for any flex or columnar table that includes a `__raw__` column.

This example illustrates querying a flex view:

1. Create a flex table.

```
=> CREATE FLEX TABLE twitter();
```

2. Load JSON data into flex table using `fjsonparser`.

```
=> COPY twitter FROM '/home/dbadmin/data/flex/tweets_10000.json' PARSER fjsonparser();
Rows Loaded
-----
10000
(1 row)
```

3. Create a flex view on top of flex table `twitter` with constraint `retweet_count > 0`.

```
=> CREATE VIEW flex_view AS SELECT __raw__ FROM twitter WHERE retweet_count::int > 0;
CREATE VIEW
```

4. Query the view. First 5 rows are displayed.

```
=> SELECT retweeted,retweet_count,source FROM (select __raw__ from flex_view) t1 limit 5;
retweeted | retweet_count |
-----+-----+-----
F          | 1              | <a href="http://blackberry.com/twitter" rel="nofollow">Twitter for
BlackBerry®</a>
F          | 1              | web
F          | 1              | <a href="http://twitter.com/download/iphone"
rel="nofollow">Twitter for iPhone</a>
F          | 23             | <a href="http://twitter.com/download/android"
rel="nofollow">Twitter for Android</a>
F          | 7              | <a href="http://twitter.com/download/iphone"
rel="nofollow">Twitter for iPhone</a>
(5 rows)
```

## Listing Flex Tables

You can determine which tables in your database are flex tables by querying the `is_flextable` column of the `v_catalog.tables` system table. For example, use a query such as the following to see all tables with a true (t) value in the `is_flextable` column:

```
=> select table_name, table_schema, is_flextable from v_catalog.tables;
```

table_name	table_schema	is_flextable
bake1	public	t
bake1_keys	public	f
del	public	t
del_keys	public	f
delicious	public	t
delicious_keys	public	f
bake	public	t
bake_keys	public	f
appLog	public	t
appLog_keys	public	f
darkdata	public	t
darkdata_keys	public	f

(12 rows)



# Setting Flex Table Parameters

Two configuration parameters affect flex table usage:

Name	Description and Use
FlexTableRawSize	<p>Determines the default column width for the <code>__raw__</code> column of a new flex table. The <code>__raw__</code> column contains the map data you load into the table. The column data type is a <code>LONG VARBINARY</code>. Setting this configuration parameter does not affect any existing flex tables. However, doing so changes the default width for any flex tables you create after changing <code>FlexTableRawSize</code>. To change an existing flex table, use the <a href="#">ALTER TABLE</a> statement, described in <a href="#">Materializing Flex Tables</a>.</p> <p><b>Default:</b> 130000</p> <p><b>Value range:</b> 1 - 32000000</p>
FlexTableDataTypeGuessMultiplier	<p>Specifies the multiplier used to set column widths when casting columns from a <code>LONG VARBINARY</code> data type for flex table views. Multiplying the longest column member by the factor pads the column width to support subsequent data loads. There is no way to determine the column width of future loads. Thus, padding adds a buffer to support values at least twice that of the previously longest value.</p> <p>These functions update the column width with each invocation:</p>

Name	Description and Use
	<ul style="list-style-type: none"><li>• <a href="#">COMPUTE_FLEXTABLE_KEYS</a></li><li>• <a href="#">COMPUTE_FLEXTABLE_KEYS_AND_BUILD_VIEW</a></li></ul> <p><b>Default:</b> 2.0: The column width multiplier. Must be a value within the following range.</p> <p><b>Range (in bytes):</b> Any value that results in a column width neither less than 20 bytes nor greater than the FlexTableRowSize value. This range is a cap to round sizes up or down, accordingly.</p>

**Note:** The FlexTableDataTypeGuessMultiplier value is not used to calculate the width of any real columns. If a flex table has defined columns, their width is set by their data type, such as 80 for a VARCHAR.

For more information, see [General Parameters](#) in the Administrator's Guide.

# Flex Data Functions Reference

The flex table data helper functions supply information you need to query the data you load. For example, suppose you don't know what keys are available in the map data. If not, you can use the [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#) function to populate a keys table and build a view. The functions aid in querying flex table and other VMap data.

Function	Description
<a href="#">COMPUTE_FLEXTABLE_KEYS</a>	Computes map keys from the map data in a <code>flextable_data</code> table, and populates the <code>flextable_data_keys</code> table with the computed keys. Use this function before building a view.
<a href="#">BUILD_FLEXTABLE_VIEW</a>	Uses the keys in the <code>flextable_data_keys</code> table to create a view definition ( <code>flextable_data_view</code> ) for the <code>flextable_data</code> table. Use this function after computing flex table keys.
<a href="#">COMPUTE_FLEXTABLE_KEYS_AND_BUILD_VIEW</a>	Performs both of the preceding functions in one call.
<a href="#">MATERIALIZE_FLEXTABLE_COLUMNS</a>	Materializes a default number of columns (50) or more or less, if specified.
<a href="#">RESTORE_FLEXTABLE_DEFAULT_KEYS_TABLE_AND_VIEW</a>	Replaces the <code>flextable_data_keys</code> table and the <code>flextable_data_view</code> , linking both the keys table and the view to the parent flex table.

While the functions are available to all users, they are applicable only to:

- Flex tables
- Associated *flex\_table\_keys* tables
- Associated *flex\_table\_view* views

By computing keys and creating views from flex table data, the functions allow you to perform `SELECT` queries. One function restores the original keys table and view that you specified when you first created the flex table.

## Flex Table Dependencies

Each flex table (*flextable*) has two dependent objects:

- *flextable\_keys*
- *flextable\_view*

While both objects are dependent on their parent table, (*flextable*), you can drop either object independently. Dropping the parent table removes both dependents, without a CASCADE option.

## Associating Flex Tables and Views

The helper functions automatically use the dependent table and view if they are internally linked with the parent table. You create both when you create the flex table. You can drop either the `_keys` table or the `_view`, and re-create objects of the same name. However, if you do so, the new objects are not internally linked with the parent flex table.

In this case, you can restore the internal links of these objects to the parent table. To do so, drop the `_keys` table and the `_view` before calling the [RESTORE\\_FLEXTABLE\\_DEFAULT\\_KEYS\\_TABLE\\_AND\\_VIEW](#) function. Calling this function re-creates either, or both, the `_keys` table and the `_view`.

The remaining helper functions perform the tasks described in this section.

## BUILD\_FLEXTABLE\_VIEW

Creates, or re-creates, a view for a default or user-defined `_keys` table, ignoring any empty keys.

### Syntax

```
build_flextable_view('flex_table' [ [, 'view_name'] [, 'user_keys_table'] ] )
```

### Arguments

<i>flex_table</i>	The flex table name. By default, this function builds or rebuilds a view for the input table with the current contents of the associated <code>flex_table_keys</code> table.
-------------------	--

<i>view_name</i>	[Optional] A custom view name. Use this option to build or rebuild a new or existing view of your choice for the input table. This option allows you to use the current contents of the associated <i>flex_table_keys</i> table, rather than the default view ( <i>flex_table_view</i> ).
<i>user_keys_table</i>	[Optional] Specifies a keys table from which to create a view. Use this option if you created a custom <i>user_keys</i> table for keys of interest from the flex table map data, rather than the default <i>flex_table_keys</i> table. The function builds a view from the keys in <i>user_keys</i> table, rather than from the <i>flex_table_keys</i> table.

## Examples

The following examples show how to call `build_flextable_view` with 1, 2, or 3 arguments.

### Creating a Default View

To create, or re-create, a default view:

1. Call the function with a single argument of a flex table, `darkdata`:

```
=> select build_flextable_view('darkdata');
        build_flextable_view
-----
The view public.darkdata_view is ready for querying
(1 row)
```

The function creates a view from the `darkdata_keys` table.

2. Query from the default view name, (`darkdata_view`):

```
=> select "user.id" from darkdata_view;
 user.id
-----
340857907
727774963
390498773
288187825
164464905
125434448
601328899
352494946
(12 rows)
```

### Creating a Custom Name View

To create, or re-create, a default view with a custom name:

1. Call the function with two arguments, a flex table, `darkdata`, and the name of the view to create, `dd_view`:

```
=> select build_flexitable_view('darkdata', 'dd_view');
        build_flexitable_view
-----
The view public.dd_view is ready for querying
(1 row)
```

2. Query from the custom view name (`dd_view`):

```
=> select "user.lang" from dd_view;
user.lang
-----
tr
en
es
en
en
it
es
en
(12 rows)
```

### Creating a View from a Custom Keys Table

To create a view from a custom `_keys` table with `build_flexitable_view`, the table must already exist. The custom table must have the same schema and table definition as the default table (`darkdata_keys`).

Create a custom keys table, using any of these three approaches: to :

1. Create a custom keys table, using any of these three approaches:
  - Create a table with the all keys from the keys table:

```
=> create table new_darkdata_keys as select * from darkdata_keys;
CREATE TABLE
```

2. ■ Alternatively, create a table based on the default keys table, but without content:

```
=> create table new_darkdata_keys as select * from darkdata_keys LIMIT 0;
CREATE TABLE
kdb=> select * from new_darkdata_keys;
  key_name | frequency | data_type_guess
-----+-----
(0 rows)
```

3. ■ Given an existing table (or creating one with no data), insert one or more keys:

```
=> create table dd_keys as select * from darkdata_keys limit 0;
CREATE TABLE
=> insert into dd_keys (key_name) values ('user.lang');
OUTPUT
-----
      1
(1 row)
=> insert into dd_keys (key_name) values ('user.name');
OUTPUT
-----
      1
(1 row)
=> select * from dd_keys;
  key_name | frequency | data_type_guess
-----+-----+-----
 user.lang |           | 
 user.name |           | 
(2 rows)
```

2. After you create your custom keys table, call the function, as shown, including:

- All arguments
- A flex table
- The name of the view to create
- The custom keys table

```
=> select build_flextable_view('darkdata', 'dd_view', 'new_darkdata_keys');
          build_flextable_view
-----
The view public.dd_view is ready for querying
(1 row)
```

3. Query the new view:

```
SELECT * from dd_view;
```

## See Also

- [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)
- [MATERIALIZED\\_FLEXTABLE\\_COLUMNS](#)
- [RESTORE\\_FLEXTABLE\\_DEFAULT\\_KEYS\\_TABLE\\_AND\\_VIEW](#)

# COMPUTE\_FLEXTABLE\_KEYS

Computes the virtual columns (keys and values) from the VMap data of a flex table and repopulates the associated `_keys` table. The keys table has the following columns:

- `key_name`
- `frequency`
- `data_type_guess`

This function sorts the keys table by `frequency` and `key_name`.

Use this function to compute keys without creating an associated table view. To also build a view, use [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#).

## Syntax

```
compute_flextable_keys('flex_table')
```

## Arguments

<i>flex_table</i>	The name of a flex table.
-------------------	---------------------------

## Examples

During execution, this function determines a data type for each virtual column. It casts the values computed to VARCHAR, LONG VARCHAR, or LONG VARBINARY. Casting choice depends on the length of the key and whether the key includes nested maps.

The following examples illustrate this function. It shows results of populating the `_keys` table, after you create a flex table (`darkdata1`) and load data:

```
=> create flex table darkdata1();
CREATE TABLE
=> copy darkdata1 from '/test/flextable/DATA/tweets_12.json' parser fjsonparser();
Rows Loaded
-----
      12
(1 row)
=> select compute_flextable_keys('darkdata1');
      compute_flextable_keys
-----
Please see public.darkdata1_keys for updated keys
(1 row)
=> select * from darkdata1_keys;
      key_name | frequency | data_type_guess
```



```

-----+-----+-----
contributors          |      8 | varchar(20)
coordinates            |      8 | varchar(20)
created_at            |      8 | varchar(60)
entities.hashtags     |      8 | long varbinary(186)
entities.urls         |      8 | long varbinary(32)
entities.user_mentions|      8 | long varbinary(674)
.
.
.
retweeted_status.user.time_zone |      1 | varchar(20)
retweeted_status.user.url      |      1 | varchar(68)
retweeted_status.user.utc_offset|      1 | varchar(20)
retweeted_status.user.verified |      1 | varchar(20)
(125 rows)

```

The flex keys table has these columns:

Column	Description
key_name	The name of the virtual column (key).
frequency	The number of times the virtual column occurs in the map.
data_type_guess	<p>The data type for each virtual column. This value is cast to VARCHAR, LONG VARCHAR or LONG VARBINARY. Casting depends on the length of the key and whether the key includes one or more nested maps.</p> <p>In the <code>_keys</code> table output, the <code>data_type_guess</code> column values are also followed by a value in parentheses, such as <code>varchar(20)</code>. The value indicates the padded width of the key column. The width is determined by calculating the longest field, multiplied by the <code>FlexTableDataTypeGuessMultiplier</code> configuration parameter value. For more information, see <a href="#">Setting Flex Table Parameters</a>.</p>

## See Also

- [BUILD\\_FLEXTABLE\\_VIEW](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)
- [MATERIALIZATE\\_FLEXTABLE\\_COLUMNS](#)
- [RESTORE\\_FLEXTABLE\\_DEFAULT\\_KEYS\\_TABLE\\_AND\\_VIEW](#)

# COMPUTE\_FLEXTABLE\_KEYS\_AND\_BUILD\_VIEW

Combines the functionality of [BUILD\\_FLEXTABLE\\_VIEW](#) and [COMPUTE\\_FLEXTABLE\\_KEYS](#) to compute virtual columns (keys) from the VMap data of a flex table and construct a view. Creating a view with this function ignores empty keys. If you don't need to perform both operations together, use one of the single-operation functions instead.

## Syntax

```
compute_flextable_keys_and_build_view('flex_table')
```

## Arguments

<i>flex_table</i>	The name of a flex table.
-------------------	---------------------------

## Examples

This example shows how to call the function for the darkdata flex table.

```
=> select compute_flextable_keys_and_build_view('darkdata');
       compute_flextable_keys_and_build_view
```

```
-----
Please see public.darkdata_keys for updated keys
The view public.darkdata_view is ready for querying
(1 row)
```

## See Also

- [BUILD\\_FLEXTABLE\\_VIEW](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [MATERIALIZER\\_FLEXTABLE\\_COLUMNS](#)
- [RESTORE\\_FLEXTABLE\\_DEFAULT\\_KEYS\\_TABLE\\_AND\\_VIEW](#)

# MATERIALIZER\_FLEXTABLE\_COLUMNS

Materializes virtual columns listed as *key\_names* in the *flextable\_keys* table you compute using either [COMPUTE\\_FLEXTABLE\\_KEYS](#) or [COMPUTE\\_FLEXTABLE\\_](#)

## KEYS\_AND\_BUILD\_VIEW.

**Note:** Each column materialized with this function counts against the data storage limit of your VerticaPremium Edition license, affecting your next Vertica license compliance audit. To manually check your Premium Edition license compliance, call the `audit()` function.

## Syntax

```
materialize_flextable_columns('flex_table' [, n-columns [, keys_
table_name] ])
```

## Arguments

<i>flex_table</i>	<p>The name of the flex table with columns to materialize. Specifying only the flex table name attempts to materialize up to 50 columns of key names in the default <i>flex_table_keys</i> table. When you use this argument, the function:</p> <ul style="list-style-type: none"><li>• Skips any columns already materialized</li><li>• Ignores any empty keys</li></ul> <p>To materialize a specific number of columns, use the optional parameter <i>n_columns</i>, described next.</p>
<i>n-columns</i>	<p>[Optional ] The number of columns to materialize. The function attempts to materialize the number of columns from the <i>flex_table_keys</i> table, skipping any columns already materialized. Vertica tables support a total of 1600 columns, which is the largest value you can specify for <i>n-columns</i>. The function orders the materialized results by frequency, descending, <i>key_name</i> when materializing the first <i>n</i> columns.</p>
<i>keys_table_name</i>	<p>[Optional] The name of a <i>flex_keys_table</i> from which to materialize columns. The function:</p> <ul style="list-style-type: none"><li>• Materializes the number of columns (value of <i>n-columns</i>) from <i>keys_table_name</i></li><li>• Skips any columns already materialized</li><li>• Orders the materialized results by frequency, descending, <i>key_name</i> when materializing the first <i>n</i> columns.</li></ul>

## Examples

The following example shows how to call `materialize_flextable_columns` to materialize columns. First, load a sample file of tweets (`tweets_10000.json`) into the flex table `twitter_r`.

After loading data and computing keys for the sample flex table, then call `materialize_flextable_columns` to materialize the first four columns:

```
=> copy twitter_r from '/home/release/KData/tweets_10000.json' parser fjsonparser();
Rows Loaded
-----
      10000
(1 row)

=> select compute_flextable_keys ('twitter_r');
      compute_flextable_keys
-----
Please see public.twitter_r_keys for updated keys
(1 row)

=> select materialize_flextable_columns('twitter_r', 4);
      materialize_flextable_columns
-----
The following columns were added to the table public.twitter_r:
      contributors
      entities.hashtags
      entities.urls
For more details, run the following query:
SELECT * FROM v_catalog.materialize_flextable_columns_results WHERE table_schema = 'public' and
table_name = 'twitter_r';

(1 row)
```

The last message in the example recommends querying the `materialize_flextable_columns_results` system table for the results of materializing the columns, as shown:

```
=> SELECT * FROM v_catalog.materialize_flextable_columns_results WHERE table_schema = 'public' and
table_name = 'twitter_r';
table_id      | table_schema | table_name |      creation_time      |      key_name      |
status |      message
-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
45035996273733172 | public      | twitter_r | 2013-11-20 17:00:27.945484-05 | contributors      |
ADDED | Added successfully
45035996273733172 | public      | twitter_r | 2013-11-20 17:00:27.94551-05 | entities.hashtags |
ADDED | Added successfully
45035996273733172 | public      | twitter_r | 2013-11-20 17:00:27.945519-05 | entities.urls      |
ADDED | Added successfully
45035996273733172 | public      | twitter_r | 2013-11-20 17:00:27.945532-05 | created_at         |
EXISTS | Column of same name already exists in table definition
(4 rows)
```

See the [MATERIALIZED\\_FLEXTABLE\\_COLUMNS\\_RESULTS](#) system table in the SQL Reference Manual.

## See Also

- [BUILD\\_FLEXTABLE\\_VIEW](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)
- [RESTORE\\_FLEXTABLE\\_DEFAULT\\_KEYS\\_TABLE\\_AND\\_VIEW](#)

## RESTORE\_FLEXTABLE\_DEFAULT\_KEYS\_TABLE\_AND\_VIEW

Restores the `_keys` table and the `_view`. The function also links the `_keystable` with its associated flex table, in cases where either table is dropped, and indicates whether one or both is restored.

## Syntax

```
restore_flextable_default_keys_table_and_view('flex_table')
```

## Arguments

<i>flex_table</i>	The name of a flex table .
-------------------	----------------------------

## Examples

This example shows how to invoke this function with an existing flex table, restoring both the `_keys` table and `_view`:

```
=> select restore_flextable_default_keys_table_and_view('darkdata');
       restore_flextable_default_keys_table_and_view
```

```
-----
The keys table public.darkdata_keys was restored successfully.
The view public.darkdata_view was restored successfully.
(1 row)
```

This example shows how you use the function to restore `darkdata_view`. However, the results indicate that `darkdata_keys` does not need restoring:

```
=> select restore_flextable_default_keys_table_and_view('darkdata');
           restore_flextable_default_keys_table_and_view
-----
The keys table public.darkdata_keys already exists and is linked to darkdata.
The view public.darkdata_view was restored successfully.
(1 row)
```

The `_keys` table has no content after it is restored:

```
=> select * from darkdata_keys;
   key_name | frequency | data_type_guess
-----+-----+-----
(0 rows)
```

## See Also

- [BUILD\\_FLEXTABLE\\_VIEW](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS](#)
- [COMPUTE\\_FLEXTABLE\\_KEYS\\_AND\\_BUILD\\_VIEW](#)
- [MATERIALIZATE\\_FLEXTABLE\\_COLUMNS](#)

# Flex Extractor Functions Reference

The following extractor scalar functions process polystructured data:

- MAPDELIMITEDEXTRACTOR
- MAPJSONEXTRACTOR
- MAPREGEXEXTRACTOR

Each function accepts input data that is:

- Existing database content
- A table
- Returned from an expression
- Entered directly

These functions do not parse data from an external file source. All functions return a single VMap value. The extractor functions can return data with NULL-specified columns.

This section describes each extractor function.

## MAPDELIMITEDEXTRACTOR

Extracts data with a delimiter character, and other optional arguments, returning a single VMap value. The `USING PARAMETERS` phrase specifies optional parameters for the function.

### Parameters

<code>delimiter</code>	VARCHAR	Single delimiter character. <b>Default value:</b>
<code>header_names</code>	VARCHAR	[Optional] Specifies header names for columns. <b>Default value:</b> <code>ucoln</code> Where <i>n</i> is the column offset number, starting with 0 for the first column. The function uses default values if you do not specify values for the <code>header_names</code> parameter.

trim	BOOLEAN	[Optional] Trims white space from header names and field values. <b>Default value:</b> true
treat_empty_val_as_null	BOOLEAN	[Optional] Specifies that empty fields become NULLs, rather than empty strings ( ' '). <b>Default value:</b> true

## Examples

These examples use a short set of delimited data:

```
Name|CITY|New city|State|zip
Tom|BOSTON|boston|MA|01
Eric|Burlington|BURLINGTON|MA|02
Jamie|cambridge|CAMBRIDGE|MA|08
```

To begin, save this data as `delim.dat`.

1. Create a flex table, `dflex`:

```
=> create flex table dflex();
CREATE TABLE
```

2. Use COPY to load the `delim.dat` file. Use the flex tables `fdelimitedparser` with the `header='false'` option:

```
=> copy dflex from '/home/release/kmm/flextables/delim.dat' parser fdelimitedparser
(header='false');
Rows Loaded
-----
                4
(1 row)
```

3. Create a columnar table, `dtab`, with an identity `id` column, a `delim` column, and a column to hold a VMap, named `vmap`:

```
=> create table dtab (id IDENTITY(1,1), delim varchar(128), vmap long varbinary(512));
CREATE TABLE
```

4. Use COPY to load the `delim.dat` file into the `dtab` table. For the `mapdelimitedextractor` function, add a header row with `USING PARAMETERS header_names=` option to specify the header row for the sample data, along with `delimiter '!'`:



```
=> copy dtab(delim, vmap as mapdelimitedextractor(delim
  USING PARAMETERS header_names='Name|CITY|New City|State|Zip')) FROM
'/home/dbadmin/data/delim.dat' DELIMITER '!';

Rows Loaded
-----
              4
(1 row)
```

5. Use `maptostring` for the flex table `dflex` to view the `__raw__` column contents. Notice the default header names in use (`ucol0 – ucol4`), since you specified `header= 'false'` when you loaded the flex table:

```
=> select maptostring(__raw__) from dflex limit 10;
      maptostring
-----
{
  "ucol0" : "Jamie",
  "ucol1" : "cambridge",
  "ucol2" : "CAMBRIDGE",
  "ucol3" : "MA",
  "ucol4" : "08"
}

{
  "ucol0" : "Name",
  "ucol1" : "CITY",
  "ucol2" : "New city",
  "ucol3" : "State",
  "ucol4" : "zip"
}

{
  "ucol0" : "Tom",
  "ucol1" : "BOSTON",
  "ucol2" : "boston",
  "ucol3" : "MA",
  "ucol4" : "01"
}

{
  "ucol0" : "Eric",
  "ucol1" : "Burlington",
  "ucol2" : "BURLINGTON",
  "ucol3" : "MA",
  "ucol4" : "02"
}

(4 rows)
```

6. Use `maptostring` again, this time with the `dtab` table's `vmap` column. Compare the results of this output to those for the flex table. Note that `maptostring` returns the `header_name` parameter values you specified when you loaded the data:

```
=> select maptostring(vmap) from dtab;
                                     maptostring
-----
{
  "CITY" : "CITY",
  "Name" : "Name",
  "New City" : "New city",
  "State" : "State",
  "Zip" : "zip"
}

{
  "CITY" : "BOSTON",
  "Name" : "Tom",
  "New City" : "boston",
  "State" : "MA",
  "Zip" : "02121"
}

{
  "CITY" : "Burlington",
  "Name" : "Eric",
  "New City" : "BURLINGTON",
  "State" : "MA",
  "Zip" : "02482"
}

{
  "CITY" : "cambridge",
  "Name" : "Jamie",
  "New City" : "CAMBRIDGE",
  "State" : "MA",
  "Zip" : "02811"
}

(4 rows)
```

7. Query the `delim` column to view the contents differently:

```
=> select delim from dtab;
      delim
-----
Name|CITY|New city|State|zip
Tom|BOSTON|boston|MA|02121
Eric|Burlington|BURLINGTON|MA|02482
Jamie|cambridge|CAMBRIDGE|MA|02811
(4 rows)
```

## See Also

- [MAPJSONEXTRACTOR](#)
- [MAPREGEXEXTRACTOR](#)

# MAPJSONEXTRACTOR

Extracts content of repeated JSON data objects, including nested maps, or data with an outer list of JSON elements. The `USING PARAMETERS` phrase specifies optional parameters for the function. Empty input does not generate a Warning or Error.

## Parameters

<code>flatten_ maps</code>	BOOLEAN	[Optional] Flattens sub-maps within the JSON data, separating map levels with a period (.). <b>Default value:</b> true
<code>flatten_ arrays</code>	BOOLEAN	[Optional] Converts lists to sub-maps with integer keys. Lists are not flattened by default. <b>Default value:</b> false
<code>reject_ on_ duplicate</code>	BOOLEAN	[Optional] Halts the load process if the file being loaded includes duplicate key names, with different case. <b>Default value:</b> false
<code>reject_ on_empty_ key</code>	BOOLEAN	[Optional] Rejects any row containing a key without a value ( <code>reject_on_empty_key=true</code> ). <b>Default value:</b> false
<code>omit_ empty_ keys</code>	BOOLEAN	[Optional] Omits any key from the load data that does not have a value ( <code>omit_empty_keys=true</code> ). <b>Default value:</b> false
<code>start_ point</code>	CHAR	[Optional] Specifies the name of a key in the JSON load data at which to begin parsing. The parser ignores all data before the <code>start_point</code> value. The parser processes data after the first instance, and up to the second, ignoring any remaining data. <b>Default value:</b> none

## Examples

These examples use the following sample JSON data:

```
{ "id": "5001", "type": "None" }  
{ "id": "5002", "type": "Glazed" }  
{ "id": "5005", "type": "Sugar" }  
{ "id": "5007", "type": "Powdered Sugar" }  
{ "id": "5004", "type": "Maple" }
```

We save this content as `bake_single.json`, and load that file.

1. Create a flex table, `flexjson`:

```
=> create flex table flexjson();  
CREATE TABLE
```

2. Use COPY to load the `bake_single.json` file with the flex tables `fjsonparser` parser:

```
=> copy flexjson from '/home/dbadmin/data/bake_single.json' parser fjsonparser();  
Rows Loaded  
-----  
5  
(1 row)
```

3. Create a columnar table, `coljson`, with an identity `id` column, a `json` column, and a column to hold a VMap, called `vmap`:

```
=> create table coljson(id IDENTITY(1,1), json varchar(128), vmap long varbinary(10000));  
CREATE TABLE
```

4. Use COPY to load the `bake_single.json` file into the `coljson` table, using the `mapjsonextractor` function:

```
=> copy coljson (json, vmap AS MapJSONExtractor(json)) FROM '/home/dbadmin/data/bake_  
single.json';  
Rows Loaded  
-----  
5  
(1 row)
```

5. Use the `maptostring` function for the flex table `flexjson` to output the `__raw__` column contents as strings:

```
=> select maptostring(__raw__) from flexjson limit 5;  
maptostring
```

```
-----  
{  
  "id" : "5001",  
  "type" : "None"  
}  
  
{  
  "id" : "5002",  
  "type" : "Glazed"  
}  
  
{  
  "id" : "5005",  
  "type" : "Sugar"  
}  
  
{  
  "id" : "5007",  
  "type" : "Powdered Sugar"  
}  
  
{  
  "id" : "5004",  
  "type" : "Maple"  
}  
  
(5 rows)
```

6. Use the `maptostring` function again, this time with the `coljson` table's `vmap` column and compare the results. The element order differs:

```
=> select maptostring(vmap) from coljson limit 5;  
      maptostring  
-----  
{  
  "id" : "5001",  
  "type" : "None"  
}  
  
{  
  "id" : "5002",  
  "type" : "Glazed"  
}  
  
{  
  "id" : "5004",  
  "type" : "Maple"  
}  
  
{  
  "id" : "5005",  
  "type" : "Sugar"  
}  
  
{  
  "id" : "5007",  
  "type" : "Powdered Sugar"
```

```
}
(5 rows)
```

## See Also

- MAPDELIMITEXTRACTOR
- MAPREGEXEXTRACTOR

# MAPREGEXEXTRACTOR

Extracts data from a regular expression and returns the results as a VMap. Use the `USING PARAMETERS pattern= phrase`, followed by the regular expression.

## Parameters

pattern=	VARCHAR	The regular expression as a string. <b>Default value:</b> An empty string ("").
use_jit	BOOLEAN	[Optional] Uses just-in-time compiling when parsing the regular expression. <b>Default value:</b> false.
record_terminator	VARCHAR	[Optional] The character used to separate input records. <b>Default value:</b> \n.
logline_column	VARCHAR	[Optional] The destination column containing the full string that the regular expression matched. <b>Default value:</b> An empty string ("").

## Examples

This example uses the following regular expression to search a Vertica log file for information that includes the `timestamp`, `date`, `thread_name`, and `thread_id`. For illustrative purposes, the regular expression components are shown here on separate lines. Remove any new line characters before using this example in your own tests.

```
^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)(?<thread_name>[A-Za-z ]+):  
(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?(?:[(?<component>\w+)\]  
<(?!<level>\w+)> )?(?<level>\w+)> @[(?<enode>\w+)\]\?: )(?!<text>.*)
```

The output in the following examples include newline characters for display purposes.

1. Create a flex table, flogs:

```
=> create flex table flogs();  
CREATE TABLE
```

2. Use COPY to load a sample log file (vertica.log), using the flex table fregexparser.

```
=> copy flogs from '/home/dbadmin/data/regdata/vertica.log' PARSER fregexparser(pattern=  
'^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d.\d+) (?<thread_name>[A-Za-z ]+):  
(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?  
(?:\[ (?<component>\w+)\] <(?<level>\w+)> )?(?:<(?<level>\w+)>  
@\[ (?<enode>\w+)\]? : )?(?<text>.*)');  
Rows Loaded  
-----  
8434  
(1 row)
```

3. Use MapToString to return the results from calling MapRegexExtractor with a regular expression. The output returns the results of the function in string format.

```
=> select maptostring(MapregexExtractor(E'2014-04-02 04:02:51.011  
TM Moveout:0x2aab9000f860-a000000002067 [Txn] <INFO>  
Begin Txn: a000000002067 \'Moveout: Tuple Mover\' using PARAMETERS  
pattern='^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d.\d+)  
(?<thread_name>[A-Za-z ]+):(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?  
(?:\[ (?<component>\w+)\] <(?<level>\w+)> )?(?:<(?<level>\w+)> @\[ (?<enode>\w+)\]? : )?  
(?<text>.*)')) FROM flogs where __identity__=13;  
  
maptostring  
-----  
----  
{  
  "component" : "Txn",  
  "level" : "INFO",  
  "text" : "Begin Txn: a000000002067 'Moveout: Tuple Mover'",  
  "thread_id" : "0x2aab9000f860",  
  "thread_name" : "TM Moveout",  
  "time" : "2014-04-02 04:02:51.011",  
  "transaction_id" : "a000000002067"  
}  
(1 row)
```

## See Also

- [MAPDELIMITEDEXTRACTOR](#)
- [MAPJSONEXTRACTOR](#)





# Flex Map Functions Reference

The flex map functions let you extract and manipulate nested map data.

The first argument of all flex map functions (except `emptymap()` and `mapaggregate()`) takes a VMap. The VMap can originate from the `__raw__` column in a flex table or be returned from a map or extraction function.

All map functions (except for `emptymap()` and `mapaggregate()`), accept either a LONG VARBINARY or a LONG VARCHAR map argument.

In the following example, the outer `maplookup()` function operates on the VMap data returned from the inner `maplookup()` function:

```
=> maplookup(maplookup(ret_map, 'batch'), 'scripts')
```

You can use flex map functions with:

- Flex tables
- Their associated `{flextable}_keys` table
- Automatically generated `{flextable}_view` views.

However, use of these functions does not apply to standard Vertica tables.

## EMPTYMAP

Constructs a new VMap with one row but without keys or data. Use this transform function to populate a map without using a flex parser. Instead, you use either from SQL queries or from map data present elsewhere in the database.

## Syntax

`emptymap()`

## Arguments

None

## Examples

### Create an Empty Map

```
=> select emptymap();
```

```

                                emptymap
-----
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
(1 row)

```

## Create an Empty Map from an Existing Flex Table

If you create an empty map from an existing flex table, the new map has the same number of rows as the table from which it was created.

This example shows the result if you create an empty map from the `darkdata` table, which has 12 rows of JSON data:

```

=> select emptymap() from darkdata;
                                emptymap
-----
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
\001\000\000\000\004\000\000\000\000\000\000\000\000\000\000\000
(12 rows)

```

## See Also

- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)

- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPAGGREGATE

Returns a LONG VARBINARY VMap with keys and value pairs supplied from two VARCHAR input columns of an existing columnar table. Using this function requires specifying an `over()` clause for the source table.

### Syntax

`mapaggregate(source_column1, source_column2)`

### Arguments

<code>source_column1</code>	Table column with values to use as the keys of the key/value pair of the returned VMap data.
<code>source_column2</code>	Table column with values to use as the values in the key/value pair of the returned VMap data.

### Examples

This example creates a columnar table `btest`, with two VARCHAR columns, named `keys` and `values`, and adds three sets of values:

```
=> create table btest(keys varchar(10), values varchar(10));
CREATE TABLE
=> copy btest from stdin;
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> one|1
>> two|2
>> three|3
>> \.
```

After populating the `btest` table, call `mapaggregate()`, using the `over` (`PARTITION BEST`) clause. This call returns the `raw_map` data:

```
=> select mapaggregate(keys, values) over(PARTITION BEST) from btest;
      raw_map
```

```
\001\000\000\000\023\000\000\000\003\000\000\000\020\000\000\000\021\000\000\000\022\000\000\000132
\003\000\000\000\020\000\000\000\023\000\000\000\030\000\000\000onethreetwo
(1 row)
```

The next example illustrates using [MAPTOSTRING\(\)](#) with the returned `raw_map` from `mapaggregate()` to see the values:

```
=> select maptostring(raw_map) from (select mapaggregate(keys, values) over(PARTITION BEST) from
btest) bit;
      maptostring
-----
{
  "one":  "1",
  "three":    "3",
  "two":  "2"
}
(1 row)
```

## See Also

- [EMPTYMAP](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPCONTAINSKEY

Determines whether a VMap contains a virtual column (key). This scalar function returns true (t), if the virtual column exists, or false (f) if it does not. Determining that a key exists before calling `maplookup()` lets you distinguish between NULL returns. The

`maplookup()` function uses for both a non-existent key and an existing key with a NULL value.

## Syntax

`mapcontainskey(VMap_data, 'virtual_column_name')`

## Arguments

<code>VMap_data</code>	Any VMap data. The VMap can exist as: <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
<code>virtual_column_name</code>	The name of the key to check.

## Examples

This example shows how to use the `mapcontainskey()` functions with `maplookup()`. View the results returned from both functions. Check whether the empty fields that `maplookup()` returns indicate a NULL value for the row (t) or no value (f):

You can use `mapcontainskey()` to determine that a key exists before calling `maplookup()`. The `maplookup()` function uses both NULL returns and existing keys with NULL values to indicate a non-existent key.

```
=> select maplookup(__raw__, 'user.location'), mapcontainskey(__raw__, 'user.location') from
darkdata order by 1;
maplookup | mapcontainskey
-----+-----
          | t
          | t
          | t
          | t
Chile     | t
Narnia    | t
Uptown..  | t
chicago  | t
          | f
          | f
          | f
          | f
(12 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPCONTAINSVALUE

Determines whether a VMap contains a specific value. Use this scalar function to return true (t), if the value exists, or false (f), if it does not.

## Syntax

```
mapcontainsvalue(VMap_data, 'virtual_column_value')
```

## Arguments

<code>VMap_data</code>	Any VMap data. The VMap can exist as: <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
<code>virtual_column_value</code>	The value whose existence you want to confirm.

## Examples

This example shows how to use `mapcontainsvalue()` to determine whether or not a virtual column contains a particular value. Create a flex table (`ftest`), and populate it with some virtual columns and values. Name both virtual columns one:

```
=> create flex table ftest();  
CREATE TABLE  
=> copy ftest from stdin parser fjsonparser();  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> {"one":1, "two":2}  
>> {"one":"one", "2":"2"}  
>> \.
```

Call `mapcontainsvalue()` on the `ftest` map data. The query returns false (f) for the first virtual column, and true (t) for the second, which contains the value one:

```
=> select mapcontainsvalue(__raw__, 'one') from ftest;  
mapcontainsvalue  
-----  
f  
t  
(2 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

# MAPITEMS

Returns information about items in a VMap. Use this transform function with one or more optional arguments to access polystructured values within the VMap data. This function requires an `over()` clause.

## Syntax

```
mapItems(VMap_data [, passthrough_arg [,...] ])
```

## Arguments

VMap_data	Any VMap data. The VMap can exist as: <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
passthrough_arg	[Optional] One or more arguments indicating keys within the map data in VMap_data

## Examples

These examples show how to use `mapItems()` in queries. All examples use `over(PARTITION BEST)`.

### Determine the Number of Virtual Columns in the Map Data

Use a flex table, labeled `darkmountain`, populated with JSON data. Query using the `count()` function. This query returns the number of virtual columns found in the map data:

```
=> select count(keys) from (select mapitems(darkmountain.__raw__) over(PARTITION BEST) from
darkmountain) as a;
count
-----
      19
(1 row)
```

### Determine What Data a Map Contains

You can also query a flex table for all of all items in the map data, as this snippet shows:

```
=> select * from (select mapitems(darkmountain.__raw__) over(PARTITION BEST) from darkmountain) as
```



```
a;
  keys      | values
-----+-----
hike_safety | 50.6
name        | Mt Washington
type        | mountain
height      | 17000
hike_safety | 12.2
name        | Denali
type        | mountain
height      | 29029
hike_safety | 34.1
name        | Everest
type        | mountain
height      | 14000
hike_safety | 22.8
name        | Kilimanjaro
type        | mountain
height      | 29029
hike_safety | 15.4
name        | Mt St Helens
type        | volcano
(19 rows)
```

### Directly Query a Key Value in a VMap

Review the following JSON input file, `simple.json`. In particular, notice the array called `three_Array`, and its four values:

```
{
  "one": "one",
  "two": 2,
  "three_Array":
  [
    "three_One",
    "three_Two",
    3,
    "three_Four"
  ],
  "four": 4,
  "five_Map":
  {
    "five_One": 51,
    "five_Two": "Fifty-two",
    "five_Three": "fifty three",
    "five_Four": 54,
    "five_Five": "5 x 5"
  },
  "six": 6
}
```

#### 1. Create a flex table, mapper:

```
=> create flex table mapper();
CREATE TABLE
```

#### 2. Load `simple.json` into the flex table mapper:

```
=> copy mapper from '/home/dbadmin/data/simple.json' parser fjsonparser (flatten_arrays=false,
flatten_maps=false);
Rows Loaded
-----
1
(1 row)
```

3. Call `mapkeys` on the flex table's `__row__` column to see the flex table's keys, but not the key submaps. The return values indicate `three_Array` as one of the virtual columns:

```
=> select mapkeys(__row__) over() from mapper;
keys
-----
five_Map
four
one
six
three_Array
two
(6 rows)
```

4. Call `mapitems` on flex table `mapper` with `three_Array` as a pass-through argument to the function. The call returns these array values:

```
=> select __identity__, mapitems(three_Array) OVER(PARTITION BY __identity__) from mapper;
__identity__ | keys | values
-----+-----+-----
1 | 0 | three_One
1 | 1 | three_Two
1 | 2 | 3
1 | 3 | three_Four
(4 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)

- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPKEYS

Returns the virtual columns (and values) present in any VMap data. This transform function requires an `over(PARTITION BEST)` clause.

### Syntax

`mapkeys(VMap_data)`

### Arguments

VMap_data	<p>Any VMap data. The VMap can exist as:</p> <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
-----------	--

## Examples

### Determine Number of Virtual Columns in Map Data

This example shows how to create a query, using an `over(PARTITION BEST)` clause with a flex table, `darkdata` to find the number of virtual column in the map data. The table is populated with JSON tweet data.

```
=> select count(keys) from (SELECT mapkeys(darkdata.__raw__) over(PARTITION BEST) from darkdata) as  
a;  
  
count  
-----  
550  
(1 row)
```

### Query Ordered List of All Virtual Columns in the Map

This example shows a snippet of the return data when you query an ordered list of all virtual columns in the map data:

```
=> select * from (SELECT mapkeys(darkdata.__raw__) over(PARTITION BEST) from darkdata) as a;  
keys  
-----  
contributors  
coordinates  
created_at  
delete.status.id  
delete.status.id_str  
delete.status.user_id  
delete.status.user_id_str  
entities.hashtags  
entities.media  
entities.urls  
entities.user_mentions  
favorited  
geo  
id  
.  
.  
.  
user.statuses_count  
user.time_zone  
user.url  
user.utc_offset  
user.verified  
(125 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)

- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPKEYSINFO

Returns virtual column information from a given map. This transform function requires an `over(PARTITION BEST)` clause.

### Syntax

```
mapkeysinfo(VMap_data)
```

### Arguments

<code>VMap_data</code>	Any VMap data. The VMap can exist as: <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
------------------------	---

### Returns

This function is a superset of the [MAPKEYS\(\)](#) function. It returns the following information about each virtual column:

Column	Description
<code>keys</code>	The virtual column names in the raw data.
<code>length</code>	The data length of the key name, which can differ from the actual string length.
<code>type_oid</code>	The OID type into which the value should be converted. Currently, the type is always 116 for a LONG VARCHAR, or 199 for a nested map that is stored as a LONG VARBINARY.
<code>row_num</code>	The number of rows in which the key was found.
<code>field_num</code>	The field number in which the key exists.

## Examples

This example shows a snippet of the return data you receive if you query an ordered list of all virtual columns in the map data:

```
=> select * from (SELECT mapkeysinfo(darkdata.__raw__) over(PARTITION BEST) from darkdata) as a;
```

keys	length	type_oid	row_num	field_num
-----+-----+-----+-----+-----				
contributors	0	116	1	0
coordinates	0	116	1	1
created_at	30	116	1	2
entities.hashtags	93	199	1	3
entities.media	772	199	1	4
entities.urls	16	199	1	5
entities.user_mentions	16	199	1	6
favorited	1	116	1	7
geo	0	116	1	8
id	18	116	1	9
id_str	18	116	1	10
.				
.				
.				
delete.status.id	18	116	11	0
delete.status.id_str	18	116	11	1
delete.status.user_id	9	116	11	2
delete.status.user_id_str	9	116	11	3
delete.status.id	18	116	12	0
delete.status.id_str	18	116	12	1
delete.status.user_id	9	116	12	2
delete.status.user_id_str	9	116	12	3

(550 rows)

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)

- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPLOOKUP

Returns values associated with a single key. This scalar function returns a LONG VARCHAR, with virtual column values, or NULL, if the virtual column does not exist. Column names are case insensitive.

Before using `maplookup()`, you can use these two functions to find out about your VMap data:

- [MAPTOSTRING](#) returns the contents of map data in a formatted text output
- [MAPCONTAINSKEY](#) determines whether a key exists in the map data

You can control the behavior for non-scalar values when loading data with the `fjsonparser` parser and its `flatten-arrays` argument. See [Loading JSON Data](#) and the [FJSONPARSER reference](#).

## Syntax

```
maplookup (VMap_data, 'virtual_column_name' [USING PARAMETERS [case_
sensitive={false | true}] [, buffer_size=n] ] )
```

## Arguments and Parameters

<code>VMap_data</code>	Any VMap data. The VMap can exist as: <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
<code>virtual_column_name</code>	The name of the virtual column whose values this function returns.
<code>buffer_size</code>	[Optional parameter] Specifies the maximum length (in bytes) of each value returned for <i>virtual_</i>

	<p><i>column_name</i>. To return all values for <i>virtual_column_name</i>, specify a <i>buffer_size</i> equal to or greater than (<math>\Rightarrow</math>) the number of bytes for any returned value. Any returned values greater in length than <i>buffer_size</i> are rejected.</p> <p><b>Default value:</b> 0 (No limit on <i>buffer_size</i>)</p>
<code>case_sensitive</code>	<p>[Optional parameter]</p> <p>Specifies whether to return values for <i>virtual_column_name</i> if keys with different cases exist.</p> <p><b>Example:</b></p> <pre>(... USING PARAMETERS case_sensitive=true)</pre> <p><b>Default value:</b> false</p>

## Examples

The following examples show ways you can use `maplookup`.

### Return Row Values of One Virtual Column

Return the row values of one virtual column, `user.location`:

```
=> select maplookup(__raw__, 'user.location') from darkdata order by 1;
maplookup
-----
Chile
Narnia
Uptown
.
.
chicago
(12 rows)
```

### Specify Buffer Size for Returned Values

Use the `buffer_size=` parameter to indicate the maximum length of any value that `maplookup` returns for the virtual column specified. If none of the returned key values can be greater than *n* bytes, you can use this parameter to allocate *n* bytes as the `buffer_size`.

In this example, simple JSON data is saved into a file, `simple_name.json`:

```
{
  "name": "sierra",
  "age": "63",
  "eyes": "brown",
```



```
    "weapon": "doggie"
  }
  {
    "name": "janis",
    "age": "10",
    "eyes": "blue",
    "weapon": "humor"
  }
  {
    "name": "ben",
    "age": "43",
    "eyes": "blue",
    "weapon": "sword"
  }
  {
    "name": "jen",
    "age": "38",
    "eyes": "green",
    "weapon": "shopping"
  }
}
```

1. Create a flex table, logs.
2. Load the simple\_name.json data into logs, using the fjsonparser. Specify the flatten\_arrays option as True:

```
=> COPY logs FROM '/home/dbadmin/data/simple_name.json' PARSER fjsonparser(flatten_
arrays=True);
```

3. Query the file to display the name key:

```
=> select name from logs ORDER BY name;
   name
-----
    ben
   janis
     jen
   sierra
(4 rows)
```

4. Use maplookup with buffer\_size=0 for the logs table name key. This query returns all of the values:

```
=> SELECT MapLookup(__raw__, 'name' USING PARAMETERS buffer_size=0) FROM logs;
   MapLookup
-----
   sierra
    ben
   janis
     jen
(4 rows)
```

5. Next, specify `buffer_size` 3, 5, and 6. Now, `maplookup()` returns only values with a byte length less than or equal to ( $\leq$ ), the specified `buffer_size`:

```
=> SELECT MapLookup(__raw__, 'name' USING PARAMETERS buffer_size=3) FROM logs;
MapLookup
-----
ben

jen
(4 rows)
=> SELECT MapLookup(__raw__, 'name' USING PARAMETERS buffer_size=5) FROM logs;
MapLookup
-----
janis
jen
ben
(4 rows)
=> SELECT MapLookup(__raw__, 'name' USING PARAMETERS buffer_size=6) FROM logs;
MapLookup
-----
sierra
janis
jen
ben
(4 rows)
```

### Interpreting Empty Fields

This example show that if you use `maplookup` without first checking the existence of a key, the output (12 empty rows) is ambiguous. When you review the following output, you cannot determine whether a `location` key:

- Exists
- Exists with a NULL value in a row
- Exists without a value

```
=> select maplookup(__raw__, 'user.location') from darkdata;
maplookup
-----
```

(12 rows)

To disambiguate empty rows, use the `mapcontainskey()` function in conjunction with `maplookup()`. When `maplookup` returns an empty field, the corresponding value from `mapcontainskey` indicates `t` for a NULL value, or `f` for no value.

The following `vsq` sample output from both functions indicates rows without values in purple. It also shows location Narnia as a value in the `user.location` virtual column:

```
=> select maplookup(__raw__, 'user.location'), mapcontainskey(__raw__, 'user.location')
from darkdata order by 1;
maplookup | mapcontainskey
-----+-----
          | t
          | t
          | t
          | t
Chile      | t
Narnia     | t
Uptown..   | t
chicago   | t
          | f >>>>>>>>No value
          | f >>>>>>>>No value
          | f >>>>>>>>No value
          | f >>>>>>>>No value
(12 rows)
```

## Query Data from Nested Maps

To access data contained in nested maps of arbitrary depth, you can call `maplookup()` recursively. The innermost map is always `__raw__`, just as a single invocation is.

The next example shows how to use `maptostring()` to return the map contents of the table `bake`, so you can see the contents of the map data:

```
=> select maptostring(__raw__) from bake;
maptostring
-----
items.item :
0.batters.batter :
  0.id : 2001
    0.type : Regular
    1.id : 2002
    1.type : Chocolate
    2.id : 2003
    2.type : Blueberry
    3.id : 2004
    3.type : Devil's Food
  0.id : 0002
  0.name : CupCake
  0.ppu : 0.55
```

```

0.topping :
0.id : 6001
    0.type : None
    1.id : 6002
    1.type : Glazed
    2.id : 6005
    2.type : Sugar
    3.id : 6007
    3.type : Powdered Sugar
    4.id : 6006
    4.type : Chocolate with Sprinkles
    5.id : 6003
    5.type : Chocolate
    6.id : 6004
    6.type : Maple
0.type : Muffin
(1 row)

```

Next, after creating a flex table and loading data, use several invocations of `maplookup()` to return values from the `map __raw__` column (). Compare the returned results to the `maptostring()` output in the previous example.

```

=> create flex table bake();
CREATE TABLE
kdb=> copy bake from '/vertica/test/flextable/DATA/bake.json' parser fjsonparser(flatten_
arrays=1,flatten_maps=0);
Rows Loaded
-----
1
(1 row)
=> select maplookup(maplookup(maplookup(maplookup(maplookup(__raw_
__, 'items'), 'item.0'), 'batters'), 'batter.0'), 'type') from bake;
maplookup
-----
Regular
(1 row)
=> select maplookup(maplookup(maplookup(maplookup(maplookup(__raw_
__, 'items'), 'item.0'), 'batters'), 'batter.1'), 'type') from bake;
maplookup
-----
Chocolate
(1 row)
=> select maplookup(maplookup(maplookup(maplookup(maplookup(__raw_
__, 'items'), 'item.0'), 'batters'), 'batter.2'), 'type') from bake;
maplookup
-----
Blueberry
(1 row)
=> select maplookup(maplookup(maplookup(maplookup(maplookup(__raw_
__, 'items'), 'item.0'), 'batters'), 'batter.3'), 'type') from bake;
maplookup
-----
Devil's Food
(1 row)

```

### Check for Case-Sensitive Virtual Columns

You can use `maplookup()` with the `case_sensitive` parameter to return results when key names with different cases exist.

1. Save the following sample content as a JSON file. This example saves the file as `repeated_key_name.json`:

```
{
    "test": "lower1"
}
{
    "TEST": "upper1"
}
{
    "TEst": "half1"
}
{
    "test": "lower2",
    "TEst": "half2"
}
{
    "TEST": "upper2",
    "TEst": "half3"
}
{
    "test": "lower3",
    "TEST": "upper3"
}
{
    "TEst": "half4",
    "test": "lower4",
    "TEST": "upper4"
}
{

```

2. Create a flex table, dupe, and load the JSON file:

```
=> create flex table dupe();
CREATE TABLE
dbt=> copy dupe from '/home/release/KData/repeated_key_name.json' parser fjsonparser();
  Rows Loaded
-----
          8
(1 row)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPSIZE

Returns the number of virtual columns present in any VMap data. Use this scalar function to determine the size of keys.

## Syntax

`mapsize(VMap_data)`

## Arguments

<code>VMap_data</code>	<p>Any VMap data. The VMap can exist as:</p> <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
------------------------	--

## Examples

This example shows the returned sizes from the number of keys in the flex table `darkmountain`:

```
=> select mapsize(__raw__) from darkmountain;
mapsize
-----
      3
      4
      4
      4
      4
(5 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPTOSTRING

Recursively builds a string representation VMap data, including nested JSON maps. Use this transform function to display the VMap contents in a readable `LONG VARCHAR` format. Use `maptostring` to see how map data is nested before querying virtual columns with `mapvalues()`.

## Syntax

```
maptostring(VMap_data [using parameters canonical_json={true  
| false}])
```

## Arguments

VMap_data	<p>Any VMap data. The VMap can exist as:</p> <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
-----------	--

## Parameters

canonical_ json	<p><i>=bool</i> [Optional parameter]</p> <p>Produces canonical JSON output by default, using the first instance of any duplicate keys in the map data.</p> <p>Use this parameter as other UDF parameters, preceded by <code>using parameters</code>, as shown in the examples. Setting this argument to <code>false</code> maintains the previous behavior of <code>maptostring()</code> and returns same-name keys and their values.</p> <p><b>Default value:</b> <code>canonical-json=true</code></p>
--------------------	---

## Examples

The following example shows how to create a sample flex table, `darkdata` and load JSON data from STDIN. By calling `maptostring()` twice with both values for the `canonical_json` parameter, you can see the different results on the flex table `__raw__` column data.

1. Create sample table:

```
=> create flex table darkdata();  
CREATE TABLE
```

2. Load sample JSON data from STDIN:



```
=> copy darkdata from stdin parser fjsonparser();
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> {"aaa": 1, "aaa": 2, "AAA": 3, "bbb": "aaa\"bbb"}
>> \.
```

3. Call `maptostring()` with its default behavior using canonical JSON output, and then review the flex table contents. The function returns the first duplicate key and its value (`"aaa": "1"`) but omits remaining duplicate keys (`"aaa": "2"`):

```
=> select maptostring (__raw__) from darkdata;
      maptostring
-----
{
  "AAA" : "3",
  "aaa" : "1",
  "bbb" : "aaa\"bbb"
}
(1 row)
```

4. Next, call `maptostring()` with using parameters `canonical_json=false`). This time, the function returns the first duplicate keys and their values:

```
=> select maptostring(__raw__ using parameters canonical_json=false) from darkdata;
      maptostring
-----
{
  "aaa": "1",
  "aaa": "2",
  "AAA": "3",
  "bbb": "aaa\"bbb"
}
(1 row)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)

- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPVALUES](#)
- [MAPVERSION](#)

## MAPVALUES

Returns a string representation of the top-level values from a VMap. This transform function requires an `over()` clause.

### Syntax

`mapvalues(VMap_data)`

### Arguments

VMap_data	<p>The VMap from which values should be returned. The VMap can exist as:</p> <ul style="list-style-type: none"><li>• The <code>__raw__</code> column of a flex table</li><li>• Data returned from a map function such as <code>maplookup()</code></li><li>• Other database content</li></ul>
-----------	--

### Examples

The following example shows how to query a `darkmountain` flex table, using an `over()` clause (in this case, the `over(PARTITION BEST)` clause) with `mapvalues()`.

```
=> select * from (select mapvalues(darkmountain.__raw__) over(PARTITION BEST) from darkmountain) as
a;
  values
-----
29029
34.1
Everest
mountain
29029
15.4
Mt St Helens
volcano
```

```
17000  
12.2  
Denali  
mountain  
14000  
22.8  
Kilimanjaro  
mountain  
50.6  
Mt Washington  
mountain  
(19 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVERSION](#)

## MAPVERSION

Returns the version or invalidity of any map data. This scalar function returns the map version (such as 1) or -1, if the map data is invalid.

## Syntax

```
mapversion(VMap_data)
```

## Arguments

VMap_data	The VMap data either from a <code>__raw__</code> column in a flex table or from the data returned from a map function such as <code>maplookup()</code> .
-----------	--

## Examples

The following example shows how to use `mapversion()` with the `darkmountainflex` table, returning `mapversion 1` for the flex table map data:

```
=> select mapversion(__raw__) from darkmountainflex;
mapversion
-----
          1
          1
          1
          1
          1
(5 rows)
```

## See Also

- [EMPTYMAP](#)
- [MAPAGGREGATE](#)
- [MAPCONTAINSKEY](#)
- [MAPCONTAINSVALUE](#)
- [MAPITEMS](#)
- [MAPKEYS](#)
- [MAPKEYSINFO](#)
- [MAPLOOKUP](#)
- [MAPSIZE](#)
- [MAPTOSTRING](#)
- [MAPVALUES](#)

# Flex Parsers Reference

Vertica supports several parsers to load different types of data into flex tables:

- [FAVROPARSER](#)
- [FCEFPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FDELIMITEDPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

Unlike with columnar tables, you must specify which parser to use when loading flex tables. You can use each flex parser to load the parser's associated type of data into columnar tables.

All parsers store the data as a single Vmap in the `LONG VARBINARY __raw__` column. If a data row is too large to fit in the column, it is rejected. Vertica supports null values for loading data with NULL-specified columns.

For information about how you can use each type of flex parser, see [Using Flex Table Parsers](#)

## FAVROPARSER

Parses data from an Avro file. The input file must use binary serialization encoding. Use this parser to load data into columnar, flex, and hybrid tables.

**Note:** The parser `favroparser` does not support Avro files with separate schema files. The Avro file must have its related schema in the file you are loading.

## Parameters

<code>flatten_arrays</code>	BOOLEAN	[Optional] Flattens all Avro arrays. Key names are concatenated with nested levels.  <b>Default value:</b> <code>false</code> (Arrays are not flattened.)
-----------------------------	---------	---

<code>flatten_maps</code>	BOOLEAN	[Optional] Flattens all Avro maps. Key names are concatenated with nested levels <b>Default value:</b> true
<code>flatten_records</code>	BOOLEAN	[Optional] Flattens all Avro records. Key names are concatenated with nested levels. <b>Default value:</b> true
<code>reject_on_materialized_type_error</code>	BOOLEAN	[Optional] Indicates whether to reject any row value for a materialized column that the parser cannot coerce into a compatible data type. See <a href="#">Using Flex Table Parsers</a> . <b>Default value:</b> false

## Examples

This example shows how to create and load a flex table with Avro data using `favroparser`. After loading the data, you can query virtual columns.

1. Create a flex table for Avro data, `avro_basic`:

```
=> CREATE FLEX TABLE avro_basic();
CREATE TABLE
```

2. Use the `favroparser` to load the data from an Avro file (`weather.avro`).

```
=> COPY avro_basic FROM '/home/dbadmin/data/flexcsv/weather.avro' PARSER favroparser();
Rows Loaded
-----
5
(1 row)
```

3. Query virtual columns from the `avro` flex table:

```
=> SELECT station, temp, time FROM avro_basic;
station | temp |      time
-----+-----+-----
mohali  | 0    | -619524000000
lucknow | 22   | -619506000000
norwich | -11  | -619484400000
ams     | 111  | -655531200000
```

```
baddi | 78 | -655509600000  
(5 rows)
```

For more information, see [Loading Avro Data](#).

## See Also

- [FCEFPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

## FCEFPARSER

Parses HPE ArcSight Common Event Format (CEF) log files. The `fcefparser` loads values directly into any table column with a column name that matches a source data key. The parser stores the data loaded into a flex table in a single VMap.

## Parameters

<code>delimiter</code>	CHAR	[Optional] Specifies a single-character delimiter. <b>Default value:</b> ' '
<code>record_terminator</code>	CHAR	[Optional] Specifies a single-character record terminator. <b>Default value:</b> <code>newline</code>
<code>trim</code>	BOOLEAN	[Optional] Trims white space from header names and key values. <b>Default value:</b> <code>true</code>
<code>reject_on_unescaped_delimiter</code>	BOOLEAN	[Optional] Determines whether to reject rows containing unescaped delimiters. The CEF standard does not permit them.

		<b>Default value:</b> false
--	--	-----------------------------

## Examples

The following example illustrates creating a sample flex table for CEF data, with two real columns, `eventId` and `priority`.

1. Create a flex table `cefdata`:

```
=> create flex table cefdata();
CREATE TABLE
```

2. Load some basic CEF data, using the flex parser `fcefparser`:

```
=> copy cefdata from stdin parser fcefparser();
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> CEF:0|ArcSight|ArcSight|2.4.1|machine:20|New alert|High|
>> \.
```

3. Use the `maptostring()` function to view the contents of your `cefdata` flex table:

```
=> select maptostring(__raw__) from cefdata;
      maptostring
-----
{
  "deviceproduct" : "ArcSight",
  "devicevendor"  : "ArcSight",
  "deviceversion" : "2.4.1",
  "name"          : "New alert",
  "severity"      : "High",
  "signatureid"   : "machine:20",
  "version"       : "0"
}

(1 row)
```

4. Select some virtual columns from the `cefdata` flex table:

```
= select deviceproduct, severity, deviceversion from cefdata;
deviceproduct | severity | deviceversion
-----+-----+-----
ArcSight      | High    | 2.4.1
(1 row)
```

For more information, see [Loading Common Event Format \(CEF\) Data](#)



## See Also

- [FAVROPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

## FCSVPARSER

Parses CSV format (comma-separated values) data. Use this parser to load CSV data into columnar, flex, and hybrid tables. All data must be encoded in Unicode [UTF-8](#) format. The parser `fcsvparser` supports the [RFC 4180](#) de facto standard for CSV data, and other options, to accommodate variations in CSV file format definitions.

The `fsvparser` does not support multibyte data. For more information about data formats, see [Checking Data Format Before or After Loading](#).

## Parameters

<code>type= {'traditional'   'rfc4180'}</code>	CHAR	[Optional] Specifies the default parameter values for the parser. You do not have to use the type parameter when loading data that conforms to the RFC 4180 standard (such as MS Excel files). See <a href="#">Loading CSV Data</a> for the RFC4180 default parameters, and other options you can specify for non-default CSV files. <b>Default value:</b> RFC4180
<code>delimiter</code>	CHAR	[Optional] Indicates the single-character value used to separate fields in the CSV data. <b>Default value:</b> ,

escape	CHAR	<p>[Optional] Specifies a single-character value. Use an escape character to interpret the next character in the data literally.</p> <p><b>Default value:</b> "</p>
enclosed_by	CHAR	<p>[Optional] Specifies a single-character value. Use and enclosed_by value to include a value that is identical to the delimiter, but should be interpreted literally. For example, if the data delimiter is a comma (,), and you want to use a comma within the data ("my name is jane, and his is jim").</p> <p><b>Default value:</b> "</p>
record_terminator	CHAR	<p>[Optional] Indicates the single-character value used to specify the end of a record.</p> <p><b>Default value:</b> \n or \r\n</p>
header	BOOLEAN	<p>[Optional] Specifies that a header column exists. If you specify true, but the data has no header, the parser uses ucol<i>n</i>, where <i>n</i> is the column offset number, starting with 0 for the first column.</p> <p><b>Default value:</b> true</p>
trim	BOOLEAN	<p>[Optional] Indicates whether to trim white space from header names and key values.</p> <p><b>Default value:</b> true</p>
omit_empty_keys	BOOLEAN	<p>[Optional] Indicates how the parser handles header keys without values. If omit_empty_keys=true, keys with an empty value in the header row are not loaded.</p> <p><b>Default value:</b> false</p>
reject_on_duplicate	BOOLEAN	<p>[Optional] Specifies whether to halt the load process if the load file includes duplicate key names with different case (such as BYPASS and BYpass).</p>

		<b>Default value:</b> false
reject_on_empty_key	BOOLEAN	[Optional] Specifies whether to reject any row containing a key without a value. <b>Default value:</b> false
reject_on_materialized_type_error	BOOLEAN	[Optional] Indicates whether to reject any materialized column value that the parser cannot coerce into a compatible data type. See <a href="#">Loading CSV Data</a> . <b>Default value:</b> false

## Examples

This example shows how you can use `fcsvparser` to load a flex table, build a view, and then query that view.

1. Create a flex table for csv data:

```
=> CREATE FLEX TABLE rfc();
CREATE TABLE
```

2. Use `fcsvparser` to load the data from STDIN. Specify that no header exists, and enter some data as shown:

```
=> COPY rfc FROM stdin PARSER fcsvparser(header='false');
Enter data to be copied followed by a newline.
End with a backslash and a period on a line by itself.
>> 10,10,20
>> 10,"10",30
>> 10,"20""5",90
>> \.
```

3. Run the `compute_flextable_keys_and_build_view` function, and query the `rfc_view`. Notice that the default `enclosed_by` character permits an escape character (") within a field ("20""5"). Thus, the resulting value was parsed correctly. Since no header existed in the input data, the function added `ucoln` for each column:

```
=> SELECT compute_flextable_keys_and_build_view('rfc');
               compute_flextable_keys_and_build_view
-----
Please see public.rfc_keys for updated keys
```

```
The view public.rfc_view is ready for querying
(1 row)

=> SELECT * FROM rfc_view;
 ucol10 | ucol11 | ucol12
-----+-----+-----
 10     | 10     | 20
 10     | 10     | 30
 10     | 20"5   | 90
(3 rows)
```

For more information and examples of using other parameters of this parser, see [Loading CSV Data](#).

## See Also

- [FAVROPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FDELIMITEDPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

## FDELIMITEDPAIRPARSER

Parses delimited data files. This parser provides a subset of the functionality in the parser `fdelimitedparser`. Use the `fdelimitedpairparser` when the data you are loading specifies pairs of column names with data in each row.

## Parameters

<code>delimiter</code>	CHAR	[Optional] Specifies a single-character delimiter. <b>Default value:</b> ' '
<code>record_terminator</code>	CHAR	[Optional] Specifies a single-character record terminator.

		<b>DDefault value:</b> newline
trim	BOOLEAN	[Optional] Trims white space from header names and key values.  <b>Default value:</b> true

## Examples

The following example illustrates creating a sample flex table for simple delimited data, with two real columns, eventId and priority.

1. Create a table:

```
=> create flex table CEFDData(eventId int default(eventId::int), priority int default
(priority::int) );
CREATE TABLE
```

2. Load a sample delimited HPE ArcSight log file into the CEFDData table, using the fcefparser:

```
=> copy CEFDData from '/home/release/kmm/flextables/sampleArcSight.txt' parser
fdelimitedpairparser();
Rows Loaded | 200
```

3. After loading the sample data file, use maptostring() to display the virtual columns in the \_\_raw\_\_ column of CEFDData:

```
=> select maptostring(__raw__) from CEFDData limit 1;
```

maptostring

```
-----
"agentassetid" : "4-WwHuD0BABCCQDVAeX21vg==",
"agentzone" : "3083",
"agt" : "265723237",
"ahost" : "svsvm0176",
"aid" : "3tGoHuD0BABCCMDVAeX21vg==",
"art" : "1099267576901",
"assetcriticality" : "0",
"at" : "snort_db",
"atz" : "America/Los_Angeles",
"av" : "5.3.0.19524.0",
"cat" : "attempted-recon",
"categorybehavior" : "/Communicate/Query",
"categorydevicegroup" : "/IDS/Network",
"categoryobject" : "/Host",
"categoryoutcome" : "/Attempt",
"categorysignificance" : "/Recon",
```

```
"categorytechnique" : "/Scan",
"categorytupledescription" : "An IDS observed a scan of a host.",
"cnt" : "1",
"cs2" : "3",
"destinationgeocountrycode" : "US",
"destinationgeolocationinfo" : "Richardson",
"destinationgeopostalcode" : "75082",
"destinationgeoregioncode" : "TX",
"destinationzone" : "3133",
"device product" : "Snort",
"device vendor" : "Snort",
"device version" : "1.8",
"deviceseverity" : "2",
"dhost" : "198.198.121.200",
"dlat" : "329913940429",
"dlong" : "-966644973754",
"dst" : "3334896072",
"dtz" : "America/Los_Angeles",
"dvchost" : "unknown:eth1",
"end" : "1364676323451",
"eventid" : "1219383333",
"fdevice product" : "Snort",
"fdevice vendor" : "Snort",
"fdevice version" : "1.8",
"fdtz" : "America/Los_Angeles",
"fdvchost" : "unknown:eth1",
"lblstring2label" : "sig_rev",
"locality" : "0",
"modelconfidence" : "0",
"mrt" : "1364675789222",
"name" : "ICMP PING NMAP",
"oagentassetid" : "4-WwHuD0BABCCQDVAeX21vg==",
"oagentzone" : "3083",
"oagt" : "265723237",
"oahost" : "svsvm0176",
"oid" : "3tGoHuD0BABCCMDVAeX21vg==",
"oat" : "snort_db",
"oatz" : "America/Los_Angeles",
"oav" : "5.3.0.19524.0",
"originator" : "0",
"priority" : "8",
"proto" : "ICMP",
"relevance" : "10",
"rt" : "1099267573000",
"severity" : "8",
"shost" : "198.198.104.10",
"signature id" : "[1:469]",
"slat" : "329913940429",
"slong" : "-966644973754",
"sourcegeocountrycode" : "US",
"sourcegeolocationinfo" : "Richardson",
"sourcegeopostalcode" : "75082",
"sourcegeoregioncode" : "TX",
"sourcezone" : "3133",
"src" : "3334891530",
"start" : "1364676323451",
"type" : "0"
```

```
}  
  
(1 row)
```

4. Select the eventID and priority real columns, along with two virtual columns, atz and destinationgeoregioncode:

```
=> select eventID, priority, atz, destinationgeoregioncode from CEFData limit 10;  
eventID | priority | atz | destinationgeoregioncode  
-----+-----+-----+-----  
1218325417 | 5 | America/Los_Angeles |  
1219383333 | 8 | America/Los_Angeles | TX  
1219533691 | 9 | America/Los_Angeles | TX  
1220034458 | 5 | America/Los_Angeles | TX  
1220034578 | 9 | America/Los_Angeles |  
1220067119 | 5 | America/Los_Angeles | TX  
1220106960 | 5 | America/Los_Angeles | TX  
1220142122 | 5 | America/Los_Angeles | TX  
1220312009 | 5 | America/Los_Angeles | TX  
1220321355 | 5 | America/Los_Angeles | CA  
(10 rows)
```

## See Also

- [FAVROPARSER](#)
- [FCEFPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

## FDELIMITEDPARSER

Parses data using a delimiter character to separate values. The `fdelimitedparser` loads delimited data, storing it in a single-value VMap. You can use this parser to load data into columnar and flex tables.

## Parameters

delimiter	CHAR	[Optional] Indicates a single-character delimiter.
-----------	------	--

		<b>Default value:</b>
record_terminator	CHAR	[Optional] Indicates a single-character record terminator. <b>Default value:</b> \n
trim	BOOLEAN	[Optional] Determines whether to trim white space from header names and key values. <b>Default value:</b> true
header	BOOLEAN	[Optional] Specifies that a header column exists. The parser uses col### for the column names if you use this parameter but no header exists. <b>Default value:</b> true
omit_empty_keys	BOOLEAN	[Optional] Indicates how the parser handles header keys without values. If omit_empty_keys=true, keys with an empty value in the headerrow are not loaded. <b>Default value:</b> false
reject_on_duplicate	BOOLEAN	[Optional] Specifies whether to halt the load process if the file being loaded includes duplicate key names with different case (such as BYPASS and BYpass). <b>Default value:</b> false
reject_on_empty_key	BOOLEAN	[Optional] Specifies whether to reject any row containing a key without a value. <b>Default value:</b> false
reject_on_materialized_type_error	BOOLEAN	[Optional] Indicates whether to reject any row value for a materialized column that the parser cannot coerce into a compatible data type. See <a href="#">Using Flex Table Parsers</a> . <b>Default value:</b> false



treat_empty_val_as_null	BOOLEAN	[Optional] Specifies that empty fields become NULLs, rather than empty strings ('').  <b>Default value:</b> true
-------------------------	---------	--

## Examples

1. Create a flex table for delimited data:

```
=> create flexible table my_test();
CREATE TABLE
```

2. Use the `fdelimitedparser` to load the data from a `.csv` file. Specify a comma (,) column delimiter:

```
=> copy my_test from '/test/vertica/DATA/a.csv' parser fdelimitedparser (delimiter=',');
Rows Loaded
-----
              3
(1 row)
```

1. Select some virtual columns from the `cefdata` flex table:

```
= select deviceproduct, severity, deviceversion from cefdata;
deviceproduct | severity | deviceversion
-----+-----+-----
ArcSight      | High    | 2.4.1
(1 row)
```

For more information, see [Loading Common Event Format \(CEF\) Data](#)

## See Also

- [FAVROPARSER](#)
- [FCEFPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FJSONPARSER](#)
- [FREGEXPARSER](#)

# FJSONPARSER

Parses and loads a JSON file. This file can contain either repeated JSON data objects (including nested maps), or an outer list of JSON elements. For a flex table, the parser stores the JSON data in a single-value VMap. For a hybrid or columnar table, the parser loads data directly in any table column with a column name that matches a key in the JSON source data.

## Parameters

flatten_maps	BOOLEAN	[Optional] Flattens sub-maps within the JSON data, separating map levels with a period (.). <b>Default value:</b> true
flatten_arrays	BOOLEAN	[Optional] Converts lists to sub-maps with integer keys. Lists are not flattened by default. <b>Default value:</b> false
reject_on_duplicate	BOOLEAN	[Optional] Halts the load process if the file being loaded includes duplicate key names, with different case. <b>Default value:</b> false
reject_on_empty_key	BOOLEAN	[Optional] Rejects any row containing a key without a value (reject_on_empty_key=true). <b>Default value:</b> false
omit_empty_keys	BOOLEAN	[Optional] Omits any key from the load data that does not have a value (omit_empty_keys=true). <b>Default value:</b> false
reject_on_materialized_type_error	BOOLEAN	[Optional] Rejects a data row that contains a materialized column value that cannot be coerced into a compatible data type ( reject_on_materialized_type_error=true). <b>Default value:</b> false
start_point	CHAR	[Optional] Specifies the name of a key in the

		<p>JSON load data at which to begin parsing. The parser ignores all data before the <code>start_point</code> value. The parser processes data after the first instance, and up to the second, ignoring any remaining data.</p> <p><b>Default value:</b> none</p>
<code>start_point_occurrence</code>	INTEGER	<p>[Optional] Indicates the <i>n</i>th occurrence of the value you specify with <code>start_point</code>. Use in conjunction with <code>start_point</code> when load data has multiple start values and you know the occurrence at which to begin parsing.</p> <p><b>Default value:</b> 1</p>
<code>suppress_nonalphanumeric_key_chars</code>	BOOLEAN	<p>[Optional] Suppresses non-alphanumeric characters in JSON key values. The parser replaces these characters with an underscore (<code>_</code>) when this parameter is true.</p> <p><b>Default value:</b> false</p>
<code>key_separator</code>	CHAR	<p>[Optional] Specifies a non-default character for the parser to use when concatenating key names.</p> <p><b>Default value:</b> ' . '</p>

## Examples

### Load JSON Data Without Optional Parameters

1. Create a flex table, `super`, with two columns, `age` and `name`:

```
=> create table super(age int, name varchar);  
CREATE TABLE
```

2. Enter values using the `fjsonparser()`, and query the results:

```
=> copy super from stdin parser fjsonparser();  
Enter data to be copied followed by a newline.  
End with a backslash and a period on a line by itself.  
>> {"age": 5, "name": "Tim"}  
>> {"age": 3}  
>> {"name": "Fred"}  
>> {"name": "Bob", "age": 10}  
>> \.
```

```
=> select * from super;  
age | name  
-----+-----  
    | Fred  
10  | Bob  
5   | Tim  
3   |  
(4 rows)
```

For other examples, see [Loading JSON Data](#).

## See Also

- [FAVROPARSER](#)
- [FCEFPARSER](#)
- [FCSVPARSER](#)
- [FDELIMITEDPARSER](#)
- [FDELIMITEDPAIRPARSER](#)
- [FREGEXPARSER](#)

## FREGEXPARSER

Parses a regular expression, matching columns to the contents of the named regular expression groups.

### Parameters

pattern	VARCHAR	Specifies the regular expression of data to match. <b>Default value:</b> Empty string ("")
use_jit	BOOLEAN	[Optional] Indicates whether to use just-in-time compiling when parsing the regular expression. <b>Default value:</b> false
record_terminator	VARCHAR	[Optional] Specifies the character used to separate input records. <b>Default value:</b> \n
logline_column	VARCHAR	[Optional] Captures the destination column containing the

		<p>full string that the regular expression matched.</p> <p><b>Default value:</b> Empty string ("")</p>
--	--	--

## Example

These examples use the following regular expression, which searches for information that includes the timestamp, date, thread\_name, and thread\_id strings. For illustrative purposes, the regular expression components are shown here on separate lines. Remove any new line characters before using this example in your own tests.

```
'^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)(?<thread_name>[A-Za-z ]+):(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?(?:\[?(?<component>\w+)\] \<(?<level>\w+)\> )?(?:<(?<level>\w+)\> @\[?(?<enode>\w+)\]\?: )?(?<text>.*).'
```

1. Create a flex table (vlog) to contain the results of a Vertica log file. For this example, we made a copy of a log file in the directory /home/dbadmin/data/vertica.log:

```
=> create flex table vlog();
CREATE TABLE
```

2. Use the fregexparser with the sample regular expression to load data from the log file (removing line characters first):

```
=> copy vlog from '/home/dbadmin/data/flex/vertica.log' PARSE FRegexParser(pattern=
'^(?<time>\d\d\d\d-\d\d-\d\d \d\d:\d\d:\d\d\.\d+)(?<thread_name>[A-Za-z ]+):
(?<thread_id>0x[0-9a-f]+)-?(?<transaction_id>[0-9a-f]+)?
(?:\[?(?<component>\w+)\] \<(?<level>\w+)\> )?(?:<(?<level>\w+)\> @\[?(?<enode>\w+)\]\?: )?
(?<text>.*).'
```

Rows Loaded

-----

8434

(1 row)

3. After successfully loading data, query columns directly from the flex table:

```
=> select text, time, thread_name from vlog limit 4;
```

text	time	thread_name
Starting up Vertica Analytic Database v7.1.0-20140331	2014-04-02 04:02:02.613	Main
Checking for missed alter partition events	2014-04-02 04:02:51.017	
ManageEpochs		
Found no missed alter partition events	2014-04-02 04:02:51.017	
ManageEpochs		
Checking for missed replace node events	2014-04-02 04:02:51.017	

ManageEpochs  
(4 rows)

4. Use the [MAPTOSTRING\(\)](#) function with the table's `__raw__` column. The four rows (`limit 4`) that the query returns are regular expression results of the `vertica.log` file, parsed with `fregexparser`:

```
=> select maptostring(__raw__) from vlog limit 4;
      maptostring
-----
{
  "component" : "Init",
  "level" : "INFO",
  "text" : "Log /scratch_b/qa/STDB/v_stdb_node0001_catalog/vertica.log opened; #2",
  "thread_id" : "0x16321430",
  "thread_name" : "Main",
  "time" : "2014-04-02 04:02:02.613"
}

{
  "component" : "Init",
  "level" : "INFO",
  "text" : "Project Codename: Dragline",
  "thread_id" : "0x16321430",
  "thread_name" : "Main",
  "time" : "2014-04-02 04:02:02.613"
}

{
  "component" : "Init",
  "level" : "INFO",
  "text" : "Processing command line: /opt/vertica/bin/vertica -D /scratch_b/qa/STDB/v_stdb_
node0001_catalog
-C STDB -n v_stdb_node0001 -h 10.20.90.192 -p 6059 -P 6061",
  "thread_id" : "0x16321430",
  "thread_name" : "Main",
  "time" : "2014-04-02 04:02:02.613"
}

{
  "component" : "Init",
  "level" : "INFO",
  "text" : "64-bit Optimized Build",
  "thread_id" : "0x16321430",
  "thread_name" : "Main",
  "time" : "2014-04-02 04:02:02.613"
}

(4 rows)
```

## See Also

- [FDELIMITEDPAIRPARSER](#)
- [FDELIMITEDPARSER](#)
- [FJSONPARSER](#)





# Send Documentation Feedback

If you have comments about this document, you can [contact the documentation team](#) by email. If an email client is configured on this system, click the link above and an email window opens with the following information in the subject line:

## **Feedback on Using Flex Tables (Vertica Analytic Database 7.2.x)**

Just add your feedback to the email and click send.

If no email client is available, copy the information above to a new message in a web mail client, and send your feedback to [vertica-docfeedback@hpe.com](mailto:vertica-docfeedback@hpe.com).

We appreciate your feedback!